

UNIVERSITY OF HASSELT

BACHELOR THESIS

---

# Machine learning techniques for flow-based network intrusion detection systems

---

*Author:*  
Axel FAES

*Advisor:*  
Prof. Dr. Peter QUAX  
Prof. Dr. Wim LAMOTTE  
*Supervisor:*  
Bram BONNE  
Pieter ROBYNS

*Bachelorproef voorgedragen tot het behalen van de graad van bachelor in de  
informatica/ICT/kennistechnologie*

*A thesis submitted in fulfillment of the requirements  
for the degree of Bachelor of Science*

*in the*

Networks and Security  
Computer Science

Academiejaar 2015-2016  
May, 2016

universiteit  
hasselt

KNOWLEDGE IN ACTION

## Declaration of Authorship

I, Axel FAES, declare that this thesis titled, “Machine learning techniques for flow-based network intrusion detection systems” and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a bachelor degree at this University.
- Where any part of this thesis has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.
- Where I have consulted the published work of others, this is always clearly attributed.
- Where I have quoted from the work of others, the source is always given. With the exception of such quotations, this thesis is entirely my own work.
- I have acknowledged all main sources of help.
- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed:

---

Date:

---

UNIVERSITY OF HASSELT

# *Abstract*

Wetenschappen  
Computer Science

Bachelor of Science

**Machine learning techniques for flow-based network intrusion  
detection systems**

by Axel FAES

---

Write the abstract

# *Acknowledgements*

---

Write the acknowledgements

# Contents

<b>Declaration of Authorship</b>	<b>i</b>
<b>Abstract</b>	<b>ii</b>
<b>Acknowledgements</b>	<b>iii</b>
<b>Nederlandstalige samenvatting</b>	<b>1</b>
0.1 Intrusion detection systems . . . . .	1
0.2 IP Flows . . . . .	1
0.3 Machine learning . . . . .	1
0.4 Conclusies . . . . .	1
<b>1 Introduction</b>	<b>2</b>
<b>2 Intrusion detection systems</b>	<b>4</b>
2.1 Host-based Intrusion Detection Systems . . . . .	5
2.2 Network-based Intrusion Detection Systems . . . . .	5
2.3 Detection . . . . .	7
2.4 Existing Intrusion detection systems . . . . .	9
<b>3 Attack Classification</b>	<b>12</b>
3.1 Classification . . . . .	12
3.2 External abnormal behaviour . . . . .	12
3.3 Internal abnormal behaviour . . . . .	14
3.4 Detection . . . . .	15
<b>4 Machine learning</b>	<b>18</b>
4.1 What is machine learning . . . . .	18
4.2 Linear Regression . . . . .	19
4.3 Classification with logistic regression . . . . .	24
4.4 Overfitting . . . . .	28
4.5 Summary . . . . .	29
<b>5 Algorithms</b>	<b>31</b>
5.1 Support Vector Machines . . . . .	31
5.2 K-Nearest Neighbors . . . . .	33
5.3 Clustering . . . . .	39
5.4 One-class Support Vector Machine . . . . .	40
5.5 Neural networks . . . . .	41
5.6 Decision tree algorithms . . . . .	44
5.7 Bayesian Algorithms . . . . .	45
5.8 Association Rule Learning Algorithms . . . . .	46
5.9 Specific algorithms for a sub-field . . . . .	46

<b>6</b>	<b>Machine Learning Techniques</b>	<b>47</b>
6.1	Dimensionality reduction . . . . .	47
6.2	Anomaly detection . . . . .	47
6.3	Online learning . . . . .	49
6.4	Bagging . . . . .	49
<b>7</b>	<b>Validating an algorithm</b>	<b>51</b>
7.1	Machine learning evaluation . . . . .	51
7.2	Error analysis . . . . .	51
<b>8</b>	<b>Machine learning for an IDS</b>	<b>56</b>
8.1	Properties of using ML for an IDS . . . . .	56
8.2	Evaluating ML for an IDS . . . . .	57
8.3	Using ML for an IDS . . . . .	58
<b>9</b>	<b>IP Flows</b>	<b>59</b>
9.1	Attributes of a flow . . . . .	59
9.2	Using IP flows . . . . .	61
<b>10</b>	<b>Implementation</b>	<b>62</b>
10.1	Technology stack . . . . .	62
10.2	Datasets . . . . .	63
10.3	Feature selection . . . . .	65
10.4	Algorithm selection . . . . .	65
<b>11</b>	<b>Evaluation</b>	<b>67</b>
11.1	Supervised learning . . . . .	67
11.2	Unsupervised learning . . . . .	67
<b>12</b>	<b>Future work</b>	<b>68</b>
12.1	Increasing accuracy . . . . .	68
12.2	Intrusion Prevention Systems . . . . .	68
12.3	Neural networks . . . . .	69
<b>13</b>	<b>Conclusion</b>	<b>70</b>
13.1	IP Flows . . . . .	70
13.2	Machine learning . . . . .	70
	<b>Bibliography</b>	<b>71</b>
<b>A</b>	<b>User guide</b>	<b>78</b>
A.1	Running . . . . .	78
A.2	Running multiple tests . . . . .	78
A.3	Config file . . . . .	78
<b>B</b>	<b>Developer Guide</b>	<b>82</b>
B.1	Machine learning module . . . . .	82
B.2	Feature module . . . . .	82
B.3	Loader module . . . . .	83
B.4	Prediction Module . . . . .	83
B.5	Training module . . . . .	83
B.6	Results module . . . . .	84
B.7	Other . . . . .	84

<b>C Meetings</b>	<b>85</b>
C.1 Meeting 1: 09 Feb 2016 . . . . .	85
C.2 Meeting 2: 12 Feb 2016 . . . . .	86
C.3 Meeting 3: 19 Feb 2016 . . . . .	86
C.4 Meeting 4: 26 Feb 2016 . . . . .	88
C.5 Meeting 5: 04 Mar 2016 . . . . .	88
C.6 Tussentijdse presentatie: 08 Mar 2016 . . . . .	89
C.7 Meeting 6: 11 Mar 2016 . . . . .	90
C.8 Meeting 7: 18 Mar 2016 . . . . .	90
C.9 Meeting 8: 24 Mar 2016 . . . . .	91
C.10 Meeting 9: 01 April 2016 . . . . .	92
C.11 Meeting Cegeka: 06 April 2016 . . . . .	93
C.12 Meeting 10: 12 April 2016 . . . . .	94
C.13 Meeting 11: 18 April 2016 . . . . .	95
C.14 Meeting 12: 22 April 2016 . . . . .	95
C.15 Meeting 13: 29 April 2016 . . . . .	96
C.16 Meeting 14: 04 Mei 2016 . . . . .	96

# List of Figures

2.1	Possible placement of IDS	4
2.2	Intrusion prevention system	6
2.3	Signature based IDS	7
2.4	Anomaly based IDS	8
2.5	The structure of the SNORT IDS	10
2.6	The structure of AI2 system	11
3.1	DDoS attack	13
3.2	Man in the middle attack	14
3.3	Botnet infection	15
4.1	Linear Regression	20
4.2	Gradient descent	21
4.3	Cost function with feature scaling	23
4.4	Cost function without feature scaling	24
4.5	Binary classification	24
4.6	Sigmoid function	25
4.7	Hypothesis function example	26
4.8	Cost function example	26
4.9	Multiclass classification	27
4.10	Overfitting	28
4.11	Effect of regularisation	29
5.1	Support Vector Machines	31
5.2	Linear and non-linear classifiers	32
5.3	K-Nearest Neighbors	33
5.4	Euclidean Distance	34
5.5	Hamming Distance	35
5.6	Chebyshev Distance	35
5.7	Manhattan Distance	36
5.8	Minkowski Distance	37
5.9	Inverse distance function	38
5.10	K-Nearest Neighbors: Uniform weights	38
5.11	K-Nearest Neighbors: Inverse distance weights	39
5.12	K-means clustering	40
5.13	One-class SVM	41
5.14	Neural network	42
5.15	Neural network with multi-class classification	42
5.16	Decision Tree Algorithm example	45
5.17	Apriori Algorithm example	46
6.1	Anomaly detection	48
6.2	Normal Distribution	49
6.3	Bagging example with Neural Networks	50



7.1	Precision and recall . . . . .	52
7.2	Training samples comparison . . . . .	55
9.1	Netflow . . . . .	59
9.2	IP API JSON format . . . . .	61
10.1	Amount of data and botnet type for each capture . . . . .	63
10.2	Distribution of labels in CTU 13 Dataset . . . . .	64

# List of Tables

7.1	Example of precision and recall of certain algorithms. . . . .	53
7.2	Example of average and F-score of certain algorithms. . . . .	53
8.1	Performance measure. . . . .	58
10.1	Labeling of the Tracelabel dataset. . . . .	65

# List of Abbreviations

<b>IDS</b>	<b>Intrusion Detection System</b>
<b>IPS</b>	<b>Intrusion Prevention System</b>
<b>IDPS</b>	<b>Intrusion Detection (and) Prevention System</b>
<b>NIDS</b>	<b>Network (based) Intrusion Detection System</b>
<b>HIDS</b>	<b>Host (based) Intrusion Detection System</b>
<b>SIEM</b>	<b>Security Information and Event Management</b>
<b>DDOS</b>	<b>Dtributed Denial of Service</b>
<b>ML</b>	<b>Machine Learning</b>
<b>MSE</b>	<b>Minimum Squared Error (function)</b>

# List of Symbols

$H_0(x)$	Hypothesis function
$J(\theta)$	Cost function

# Nederlandstalige samenvatting

---

Write the samenvatting in dutch

**0.1 Intrusion detection systems**

**0.2 IP Flows**

**0.3 Machine learning**

**0.4 Conclusies**

# Chapter 1

## Introduction

The internet is constantly growing and new network services arise constantly. Sensitive data is also increasingly being stored digitally. All these new services could contain security flaws which could leak private data, such as passwords or other sensitive data. This means that security flaws become more and more important since they can cause so much damage. It is not just the leaking of sensitive data that is an issue, but also protecting a computer or network against malware is important.

For example, earlier this year, the internal network of a hospital was attacked by a ransomware attack. The attackers did not get access to personal information and the hospital was able to go back to doing paperwork. But it does show that cyber attacks are a serious threat. [1]

Considering this, it becomes more important to be able to detect and prevent attacks on network systems. Intrusion detection systems are used for this purpose. An intrusion detection system can alert administrators of malicious behaviour. Intrusion detection systems can detect attacks using several different methods as explained in Section 2.

These methods are not always reliable. They may not catch every attack in which case the intrusion detection system is not that useful. They could also give alerts for events that are not malicious at all. This could desensitize administrators for the alerts an intrusion detection system gives.

In order to have good performance, most intrusion detection systems need a lot of manual maintenance. This is needed to constantly check the alerts the intrusion detection system gives and tweak the system if required. Other systems which are already well configured have a big price tag. This thesis tries to find out *whether an intrusion detection system can work out-of-the-box with an acceptable performance*.

This is done by using machine learning algorithms. These are algorithms which can learn and find patterns in input. They can find out what different types of behaviour look like. This means that they could for example learn what botnet behaviour looks like. Machine learning algorithms are further explained in Section 4. *Machine learning algorithms seems promising for the problem of automatic intrusion detection, but is this true?* However, there are different kinds of machine learning algorithms. *These different kinds of algorithms will be discussed and evaluated to see whether they are useful in anomaly detection.*

The fact that the internet is growing continuously, this means that more and more data passes through the networks of big companies and data centers. This makes it quite difficult to be able to check every piece of data that passes through the network. This could be solved by using IP Flows. These aggregate a lot of packet data into a small flow, they are further explained in Section 9. These flows are small, but also contain less information than the complete packet data. *But is it possible to still detect attacks when only the IP flows are used?*

It is not just important to find out whether attacks can still be detected when only IP flows are used. It is also important to know which kind of anomalies can be detected. *Could every type of attack be detected or is it a subset?*

These questions lead to the development of a system which uses machine learning techniques in a purely flow based intrusion detection system. Such a system would be one of the first lines of defense of a company against cyber attacks. Real world data is going to be used to test *whether the system can operate in a real life situation.*

## Chapter 2

# Intrusion detection systems

An intrusion detection system is a system which tries to determine whether a system is under attack, to detect intrusions within a system. Intrusion detection systems are often called IDS's. Intrusions can also be called attacks or anomalies. It does this by monitoring network or system activities. One way of categorizing IDS's is based on the method of detection intrusion. The first type would be monitoring the network, these are called network-based intrusion detection systems, or NIDS. When the intrusion detection system only monitors system activities, it is called a host-based intrusion detection system, or HIDS. [2]

Figure 2.1 shows the possible placements of an IDS. It can be placed before any firewall, being the first defense of a network. This is a NIDS. An IDS can also be placed within a network. This IDS can still monitor the network but it can also monitor system activities of the workstations. An HIDS is most commonly, but not always located on the device that it monitors.

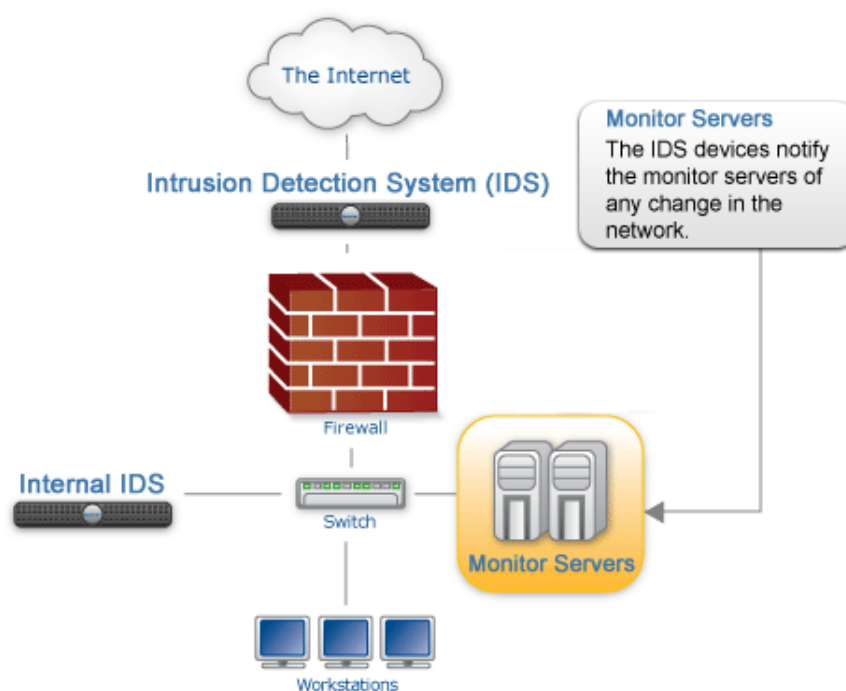


FIGURE 2.1: An IDS can for example be placed within the network or just before the network. [3]



## 2.1 Host-based Intrusion Detection Systems

Host-based intrusion detection systems are systems that monitor the device on which they are installed, or directly connected to. The way they monitor the system can range from monitoring the state of the main system through audit logs, to monitoring program execution. In this way they can be quite indistinguishable from Anti-Virus programs. [2]

Since HIDS rely so much on audit logs, they can become limited by them. If the software that is being monitored does not provide enough information in audit logs, the IDS cannot always determine intrusions successfully. Another issue can be the sheer volume of the audit logs. Every monitored log needs to be parsed, this means that the HIDS can have a big impact on the performance of the host system if it is installed there. [4]

Another disadvantage is that any vulnerability that causes the audit files to be changed, also impacts the integrity of the HIDS. If an audit file is changed, the HIDS cannot see and detect what truly happened. [5]

## 2.2 Network-based Intrusion Detection Systems

Network-based intrusion detection systems are placed at certain points within a network in order to monitor traffic from and to devices within the network. They operate on the same concept as wiretapping. They "tap" into a network and listen to all communication that happens. [2]

Using the network data instead of audit trails such as HIDS is desirable in multiple ways. One advantage is that they do not impact the performance of programs that are using the network. It is also more difficult for an intruder to attack the IDS itself. Network traffic is always visible and cannot be changed like an audit trail. The intruder could try to minimize his network activity, but the risk is lower. NIDS are also more portable than HIDS. They monitor traffic over a network and are independent of the operating system they run on. [4]

The system can analyse the traffic using multiple techniques to determine whether the data is malicious. There are two different ways to analyse the network data. The analysis can be packet-based or flow-based.

Packet-based analysis uses the entire packet including the headers and payload. An intrusion detection system that uses packet-based analysis is called a packet-based network intrusion detection system. The advantage of this type of analysis is that there is a lot of data to work with. Every single byte of the packet could be used to determine whether the packet is malicious or not. The disadvantage is immediately obvious once we look at networks through which a lot of data passes, such as data centers. Analysing every byte is very work-intensive and near impossible to do in such environments. [6]

Flow-based analysis doesn't use individual packets but uses general aggregated data about network flows. An intrusion detection system that uses flow-based analysis is called a flow-based network intrusion detection system. A flow is defined as a single connection between the host and another device. A flow can be defined using a (source\_IP, destination\_IP, source\_port, destination\_port) tuple. However flows also contains other information. IP Flows are discussed in depth in Section 9. Because of the other information, flows can still contain a lot of information, even when compared to packets. Since flow data is much more compact than all the individual packets, it is much more feasible for data centers to use flow-based intrusion detection systems. [6] [7]

### 2.2.1 Intrusion Prevention Systems

An intrusion prevention system or IPS/IDPS is an intrusion detection system that also has the ability to prevent attacks. An IDS does not necessarily need to be able to detect attacks at the exact moment they occur, although it is preferred. An IPS needs to be able to detect attacks real-time since it also needs to be able to prevent these attacks. For network attacks these prevention actions could be closing the connection, blocking an IP or limiting the data throughput. [8]

The change to requiring attacks to be detected at real time can severely impact the methods that are used to detect these attacks. For example, an IDS might give an alert even though the IDS is not certain that whatever it is alerting is actually an anomaly. An IPS needs to be certain before it can take action. Otherwise the IPS might take actions which the business employing the IPS does not want. [8]

In case of a NIPS, an network-based intrusion prevention system, some advantages of being network-based are not true anymore. An NIPS needs to see all data, and preferably block network data before it reaches its destination as seen in Figure 2.2. This means that the performance of programs using the network might be affected by the NIPS. [9]




FIGURE 2.2: An Intrusion prevention system. [10]

## 2.3 Detection

There are multiple different methods to detect intrusions. There are **Signature based methods** and there are **Anomaly Based** methods. Both of these methods have their own strengths and weaknesses.

### 2.3.1 Signature based methods

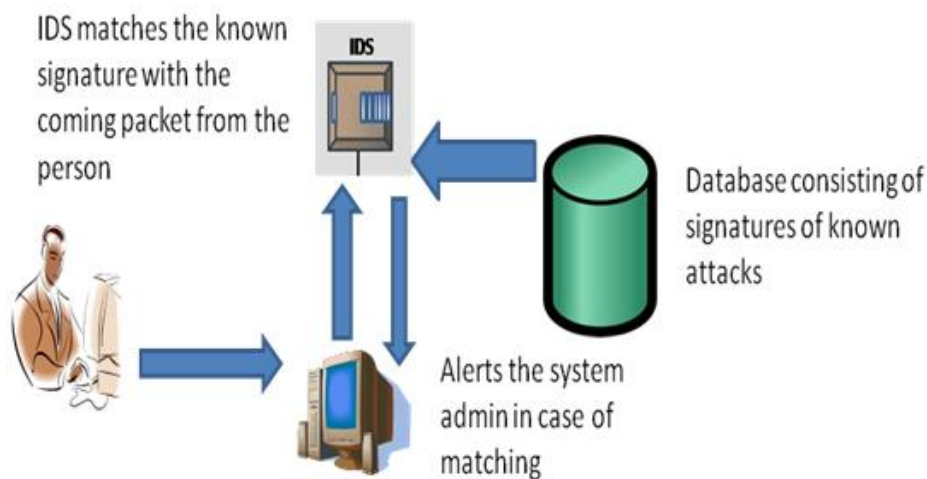


FIGURE 2.3: An Signature-based intrusion detection system. [11]

Signature based methods compare so called "signatures" with an existing database of signatures. A packet or flow record is decomposed into features that together construct a signature. If the signature of an incoming flow or packet matches with a signature in the database, it is flagged as malicious. Pseudocode of this can be seen below. Signature-based methods have little overhead in both computation and preprocessing as it only tries to match incoming signatures to known signatures in the database. Because it only compares signatures, it is easy to deploy within a network. The system does not need to learn what the traffic within a network looks like. [12]

---

```

signatures = get_signatures_from_database()
while True:
    packet = get_next_packet()
    packet_signature = get_important_features(packet)

    if (signatures.contains(packet_signature):
        generate_alert(packet)
    else:
        # packet signature was not
        # a match with known signatures
        continue

```

---

Signature based methods are very effective against known attacks. New

attacks cannot be detected unless the database is updated with new signatures. It is also possible for attackers to avoid being caught by signature based methods, only a slight modification of the "signature" is required in order to bypass the exact matching. [13]

This could be done by trying to make the network behaviour look more like normal behaviour. For example, a botnet uses IRC communication with his master. The communication might follow certain patterns which are different from usual IRC traffic from a chat. The creator might change the botnet so that communication between the botnet and the master look similar to the usual IRC traffic. This change causes the network behaviour to be different from the signature database and not generate an alert. Updating the signature database requires a lot of technical effort, since new attacks are discovered all the time.[12] [14]

### 2.3.2 Anomaly based methods




FIGURE 2.4: An Anomaly-based intrusion detection system. [11]

Anomaly based methods, also called Behaviour based methods are methods in which the IDS tries to model the behaviour of network traffic. When an incoming packet deviates from this model, it is flagged as malicious and an alert is sent. Because they use a statistical model of normal behaviour, they should be able to detect all deviations from this normal behaviour. As a result, new attacks that deviate too much from normal behaviour are detected as well. [11]

Since a model of the network traffic needs to be created, the system cannot be deployed into a network and be expected to work. The system needs to learn the behaviour of the network traffic. Problems, such as generating a lot of false positive alarms, can arise when training data includes mistakes, such as misclassifications.

Machine learning algorithms can be used as an anomaly based method. Machine learning techniques have the ability to learn from data and decide whether new data is malicious. [11]

## 2.4 Existing Intrusion detection systems

There are a lot of different intrusion detection and intrusion prevention systems on the market. Some of these systems are very expensive, others are completely open source and free.

### 2.4.1 Alienvault

*Alienvault* is a business that develops software to manage cyber attacks. They create SIEM solutions, Security Information and Event Management solutions. These are tools that provide methods to analyse security threats. IDS's are incorporated into SIEM's. [15]

The product they make is the *AlienVault Unified Security Management* (USM). USM is an all-in-one tool. USM contains a lot of tools that are required to analyse a system for cyber security threats. It contains tools that can scan and test the network for vulnerabilities.

USM supports both host-based intrusion detection and network-based intrusion detection. It also has the option to do both full packet and netflow analysis, supporting both types of NIDS systems. [16]

*Alienvault* works on both signature-based and anomaly-based intrusion detection. However, it is interesting to note that they are working on new strategies that can detect intrusions. Research is being done using neural networks to make intrusion detection more accurate. [17]

### 2.4.2 SNORT

Snort is an open-source network intrusion detection and prevention system. It is made to be very lightweight in use. It runs on UNIX derivatives and Windows. SNORT is the most widely used IDS worldwide and has become the in reality standard for the industry. SNORT can also operate on both packet level and IP flow level. However, SNORT uses a combination of signature and anomaly detection methods. [18]

SNORT uses a system based on rules. Using rules, an administrator can configure SNORT. These rules allow for both signature and anomaly detection methods. This means that SNORT itself does not work with machine learning algorithms. However, the rules can be made using machine learning algorithms. These algorithms make rules based on which traffic is classified as normal and abnormal. [19]

SNORT can, since a couple years, also be deployed as a intrusion prevention system. This is done by adding a new type of rule, a "drop" rule. A "drop" rule has higher precedence than an "alert" rule. This means that any packet that matches a "drop" rule and an "alert" rule will be dropped.

Since SNORT is open source, the internal structure can be observed. There are several components that work together to detect abnormal behaviour and to generate output in a format that is appropriate for an intrusion detection system. [18]

- Packet sniffer and decoder
- Preprocessors
- Intrusion detection engine
- Logger and alert system
- Output modules

In Figure 2.5, the structure of the SNORT components can be seen. The **packet sniffer and decoder** sniffs packets from different network interfaces. The packets are then prepared to go to the preprocessor. The **preprocessor** can already extract the most important data from packets and might even modify the data. In the **detection engine** rules are used to detect any intrusions. The rules are matched against every packet. If a packet is matched, it may be dropped or an alert can be generated, depending on the type of the rule.

Alerts can be logged to different kinds of files. This happens in the **logger and alert system**. These could be text files or tcpdumps. The **output modules** can do different operations on the output to generate new log files. [18]

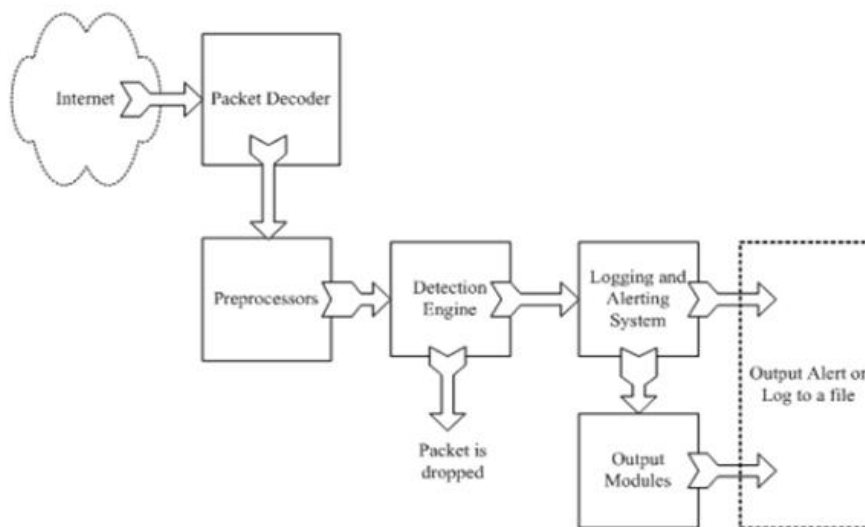


FIGURE 2.5: The structure of the SNORT IDS. [11]

### 2.4.3 MIT AI<sup>2</sup>

MIT is also working on methods to use machine learning to defend against cyber attacks. In their paper "AI<sup>2</sup>: Training a big data machine to defend", they

present a new method. Their system has four components. A big data processing system, an outlier detection engine, a mechanism to obtain feedback from security analysts, and a supervised learning module. [20]

Their system tries to combine the expertise of security experts, and the speed and ability to detect new attacks of machine learning. More specifically, they use unsupervised machine learning. They preferred to use unsupervised machine learning since labeled data is rare and attacks constantly evolve. In the system they generate their own labels and use a supervised learning algorithm with these labels.

The **big data processing system** is a system that can extract features of different entities from raw data. The **outlier detection engine** is a system that uses unsupervised learning. It uses the features that have been found in the big data processing system. They use three different methods, density, matrix decomposition, or replicator neural networks. [20]

The output of this unsupervised system is processed and shown to a **security analyst**. The security analyst can verify or refute the output. The feedback is fed to a **supervised learning algorithm**. The supervised learning algorithm learn a model that can use this feedback to better predict whether any new event is normal or abnormal. With more feedback, the system becomes more and more correct. The flow of the system can be seen in Figure 2.6. [20]

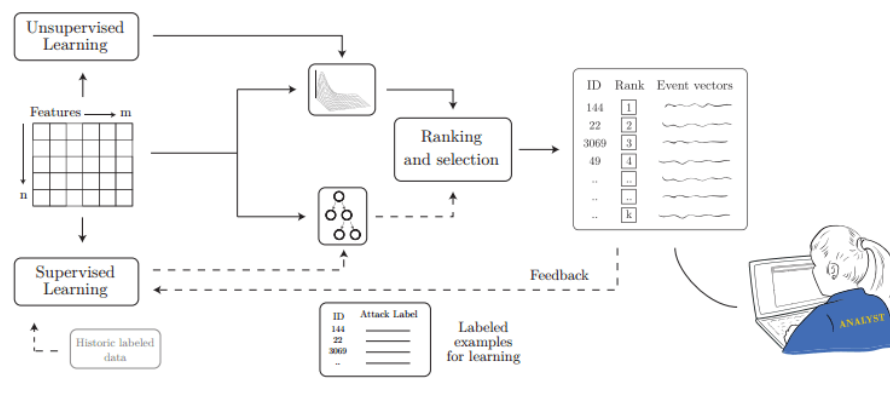


FIGURE 2.6: The structure of AI2 system. [21]

Their system has been tested by monitoring a large web-scale platform. This platform generated millions of log lines per day. The monitoring lasted three month and generated 3.6 billion log lines. They could detect 85 percent of attacks and reduced the amount of false positives by 5 times since the previous implementation. [21]



## Chapter 3

# Attack Classification

An intrusion detection system can use multiple methods to detect malicious behaviour. In order to make the IDS as effective as possible, the exact classifications of malicious behaviour that can be detected need to be known. Classification can happen on several different ways. For example they can be classified using similar behaviour or they could be classified using similar damage caused. [22]

### 3.1 Classification

A usefull classification is to first make a distinction between internal and external malicious behaviour. This makes it easier for humans to understand. The IDS itself can work with different kind of classifications. However, the IDS has to communicate with an administrator about the detections. A distinction between internal and external malicious behaviour is easier to understand.

The exact classifications are not mutually exclusive. Some types of malicious behaviour can be both internal and external. The exact classification needs to be much deeper than just internal and external. Every type of malicious behaviour is identified by different characteristics. Knowing these characteristics is useful to be able to tweak the IDS to make identification more effective. [23]

### 3.2 External abnormal behaviour

External abnormal behaviour consists of different kind of attacks on a systems. There are many different type of attacks. There are **Physical attacks**, **Buffer overflows**, **Distributed Denial of Service**, **Brute-force attacks**, **Vulnerability scans** and **Man in the middle attacks**.

#### 3.2.1 DDoS

**DDOS** or Distributed Denial of Service attacks are attacks which attempt to make a network resource temporarily or permanently unavailable for the users of that resource. An attack could happen by flooding a system with TCP SYN packets. In Figure 3.1 an example can be seen. The attacker uses some Master Nodes to control a number of Daemon Nodes. These Daemon Nodes are the computers that actually carry out the attack.



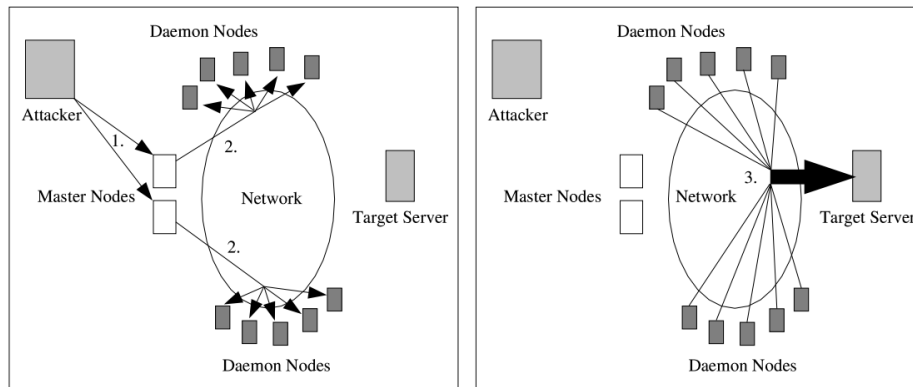


FIGURE 3.1: DDoS attack. [23]

### 3.2.2 Vulnerability scans

**Network scans** are information gathering attacks. They do not cause any damage by themselves but usually serve the purpose to gather information about a system that could be used in further attacks. Network traffic sniffing or port scans are examples of network scans. [13]

### 3.2.3 Buffer overflows

Buffer overflows are one of the most used methods of attacking a computer or network. They are used to exploit flawed programming. They work by trying to overflow buffers within the program. When this happens the overflowed data will overflow into adjacent memory. This can cause data corruption or alter the execution of a program. There are stack-based and heap-based buffer overflows. Stack-based overflows try to overflow the program stack and heap-based overflows try to overflow dynamically created memory in a program. [22]

### 3.2.4 Brute-force attacks

Brute-force attacks are a method that can be used to gain access to a target computer. This can be done by for example attempting all different kind of passwords. These attacks will eventually succeed but they can take a long time to execute. [22]

### 3.2.5 Physical attacks

Physical attacks are attacks that try to physically damage computers. These attacks can be very basic, such as cutting a wire. However, they can also be energy attacks, such as electro-magnetic pulses (EMP's) or high and low energy radio frequency (HERF and LERF) attacks. [23]

### 3.2.6 Man in the Middle

A man in the middle attack (MITM) is an attack where the attacker tries to secretly relay and possibly alter communication between two computers or networks as can be seen in Figure 3.2.




FIGURE 3.2: Man in the middle attack. [24]

### 3.3 Internal abnormal behaviour

Internal abnormal behaviour can be called malware. There are several types of malware. There are four distinct categories of malware. There are **bot-nets**, **viruses**, **trojan horses** and **worms**. Malware are actual programs that infect a system to execute a specific task. The task of the malware defines which category the malware belongs in.

#### 3.3.1 Trojan horses

**Trojan horses** are programs disguised as harmless applications but contain malicious code. They get their name from the battle of Troy. In stories it is said that the Greeks created a giant wooden horse and filled it with soldiers. The Trojans thought the horse was a gift of surrender. They brought the horse inside Troy. At night, the Greeks came out the horse, opened the gates of Troy and defeated Troy. This is also the way a Trojan horse works. Some of the capabilities of a Trojan horse include keystroke logging, program installation and file transfers. [23]

#### 3.3.2 Viruses

**Viruses** are self-replicating programs that infect and spread through other programs. An infection means that they embed themselves within another program. When that program is run, the virus also activates and runs. [23]

#### 3.3.3 Worms

**Worms** are programs that replicate themselves among a network. They can spread extremely fast. Unlike viruses, they do not require infected files to spread. [23]

### 3.3.4 Botnets

**Botnets** is malware that causes infected computers to become "slaves" to the master. An infected computer is controlled externally by the bot-master without the knowledge of the owner of the infected computer. An example of this is shown in Figure 3.3. The bot-master can use the distributed network of "slave" computer to perform other malicious tasks, such as performing an DDOS attack. [13]




FIGURE 3.3: Botnet infection. [25]

## 3.4 Detection

An NIDS only monitors the network. As such not every attack can be detected by an NIDS. Only the attacks that actually use the network can be detected. Flow-based IDS have the additional constraint that they can only use flow data. This further limits the attacks that can be detected. The attacks that can generally be detected using a flow-based network intrusion detection systems are:

- DDOS
- Vulnerability scans
- Worms
- Botnets

Other attacks either do not use network communication, or they are not visible within the header information of network traffic. In order to detect other attacks, including **Viruses**, **Trojan horses**, **Man in the Middle** and **Buffer overflows**, other detection systems such as HIDS or Packet-based NIDS should be used. **Physical attacks** are even more difficult to detect.

However there are some special circumstances in which other types of abnormal behaviour can be detected. When more information about a network and the abnormal behaviour is known, other types of abnormal behaviour can be detected. For example, when it is known that viruses are always distributed from the same source, it is possible to detect this distribution. This is analogous for other attacks.

Other activities such as phoning home which a keystroke logging Trojan might do, could also be detected. **Brute-force attacks** can also be detected since they might generate a lot of network traffic. However, unless the **Brute-force attack** is also a DDoS attack, the traffic mostly looks harmless from a flow-based approach. [22]

### 3.4.1 Distributed Denial of Service

A distributed denial of service can be detected by the amount of data that is being received. However, there are many different types of DDoS attacks. There are ICMP floods, SYN floods, etc. These attacks can be described in terms of traffic patterns. A traffic pattern is expressed in a couple features. These features include the number of flows and packets, the packet size, and the total bandwidth used during the traffic. For example UDP flooding can be characterised by a traffic pattern which contains a lot of packets. These patterns can be searched for during the detection phase. [26]

### 3.4.2 Network scans

There are three categories of network scans.

- Horizontal scans: a single port is scanned across many different devices.
- Vertical scan: several different ports are scanned on a single device
- Block scan: a combination of both a vertical and a horizontal scan.

Scans can also be described using traffic patterns. They are characterised with a high number of flow and a low number of packets. These can again be used to detect whether a vertical or horizontal scan occurs. [13] [26]

### 3.4.3 Worms

Worms exhibit different behaviour depending on their current state. There are two different states, a target discovery state and a transfer state. In the target discovery state, the worm explores the network to find vulnerabilities and a host to infect. During the transfer state, the worm actually transfers itself to the targeted host. The Sapphire/Slammer worm is an example of this type of behaviour. [27]

Since transferring of the worm itself happens within the payload data, a flow-based NIDS cannot detect this state. The target discovery state can be detected. Worms use techniques similar to network scans in order to find vulnerable hosts. So similar detection techniques can be used to detect worms. [28]

#### 3.4.4 Botnets

Botnets usually consist out of a huge amount of infected slaves controlled by a central bot-master. Locating the individual infected slaves and isolating them is a difficult problem but is also insignificant due to the huge amount of remaining slaves. Detecting the bot-master and isolating that device is key to taking down a botnet. However indentifying botnet behaviour is a far more difficult problem than detection other types of malicious activities. [29] Malicious behaviour alone is not enough to detect botnets.

Botnets often use IRC channels in order to communicate between slaves and the bot-master. These can be indentified using flows since they often use specific ports. It is possible to use a method that does not require specific port numbers. This requires flows including extra information such as the number of packets for which the PUSH flag is set.

## Chapter 4

# Machine learning

### 4.1 What is machine learning

Machine learning is a subfield of Computer Science. It is a type of Artificial Intelligence which allows programs to learn and find patterns within data. [30]

Machine learning explores algorithms that can learn from and make predictions on data. These algorithms are called machine learning algorithms. A machine learning algorithm has to learn before it can be used to make predictions on data. **Learning** means that the algorithm has to be shown several examples of data and what the correct predictions for these examples would be. The amount of examples that have to be shown to the algorithm can be within the range of several thousands.

Once the machine learning algorithm has learned from the data, it can be used to make predictions on other data. For example, machine learning can be used in order to watch the heartrate of patients at a hospital. During the learning phase, the machine learning algorithm is shown the heartrate of a patient and the current time. After learning is done, the machine learning algorithm can predict what the heartrate of that patient should be based on the current time. This can be used to determine whether the patients heartrate is normal by comparing the predicted heartrate and the real heartrate.

There are two classes of machine learning algorithms. There is **supervised** learning and **unsupervised** learning. Supervised learning is trained using labeled data. Labeled data is data which consists of input data and the corresponding output data. Unsupervised learning uses unlabeled data. The data used to train machine learning algorithms is called a **training set**.

Supervised learning is the type of learning when the data presented to the learning algorithm is labeled. This means that the user of the learning algorithm can already make distinctions within the presented data. This is useful for regression or classification problems. Regression is explained in Section 4.2 and Classification is explained in Section 4.3.

Unsupervised learning is the type of learning when the data presented to the learning algorithm is unlabeled. The learning algorithm has to find structure in the given data.

A **training sample** is a data point in an available training set that is used in a predictive modeling task. For example, if machine learning is applied to spam filtering, the training set is a collection of emails, some of which are spam emails. A training sample would be a single email. Alternative names are a training example or training instance.

Machine learning is applied for predictive modeling. In predictive modeling, a particular process or event is modeled. Using a training set, the model tries to learn or approximate a particular function that, for example, lets us distinguish spam from non-spam email. This function is called the target function. The function the model makes to approximate the target function is called the **hypothesis**.

To model the particular process or event, one or more **features** are extracted. A feature is an individual property. All features combined define the model of the particular process or event. In the example of mails, a feature could be the textbody, or it could be the sender of the email. It could even be individual characters which appear in the email. Which features are chosen is completely dependent on which problem is being solved.

This chapter gives an introduction to the mathematical background of supervised machine learning. It starts with linear regression which is a simple category of machine learning algorithms. The principles behind linear regression are the same as for other machine learning algorithms. Afterwards logistic regression is used to explain how classification works within machine learning. It begins by explaining what exactly the hypothesis is and how it can be constructed. [31]

## 4.2 Linear Regression

Linear regression is a statistical approach to model the relationship between an "output" value  $y$  and one or more "input" values  $x_1 \dots x_n$ . It belongs to the category of supervised learning. An example of this can be seen in Figure 4.1. The black dots represent the data to be modeled. The blue line is the model. This example only has one "input" value. This specific case of regression is called simple linear regression. [32].

### 4.2.1 Hypothesis

Machine learning relies heavily on a hypothesis. This is a function that transforms a given input to the machine learning algorithm into the required output. It is a function that tries to model the target function, it tries to give a function which fits the input data. When only one feature is considered, a hypothesis is a function of the form [33]:

$$H_0(x) = \theta_0 + \theta_1 * x \quad (4.1)$$

For example, we could use the grades of high school students to predict their chance of success at university. The  $x - axis$  represents the grades (on a scale from 0 to 10) for students and the  $y - axis$  is the chance of success. Given is input data (the black points) and from that data a hypothesis (the

blue line) is constructed.

The hypothesis function can be generalised for  $N$  features. It generally has the form:

$$H_0(x) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n \quad (4.2)$$

$x_0$  always has the value 1. The  $\theta$  values and the  $x$  values can be represented using vectors. Now the hypothesis function can be written in a more convenient way:

$$H_0(x) = \theta^T X = [\theta_1, \theta_2, \dots, \theta_n]^T [x_0, x_1, \dots, x_n] \quad (4.3)$$

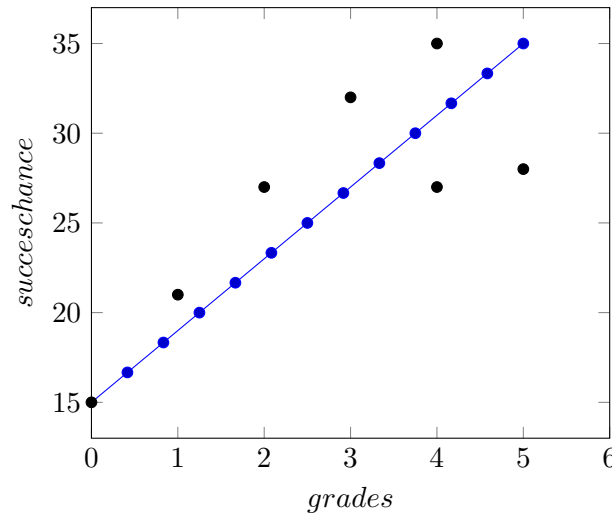


FIGURE 4.1: Linear Regression

The hypothesis is the core of a machine learning algorithm. During the training process, the algorithm learns what the hypothesis looks like. More specifically, it tries to find correct values for the vector  $\theta$  (as seen in equation 4.3). Once the hypothesis has been determined, the difficult work is done. The hypothesis forms the predictive model. The prediction itself is straightforward. A new datapoint is entered into the hypothesis equation and it returns the predicted value.

#### 4.2.2 Cost function

In order to construct a good hypothesis function, good values of  $\theta$  have to be chosen. This can be seen as a minimization problem. The difference between any output  $y$  and  $H_0(x)$  has to be minimized. More concretely, the squared difference has to be minimized. This is the MSE, "Minimum Squared Error" function or also called the cost function. The notation  $x^{(i)}$  and  $y^{(i)}$  is to denote the  $i$ th training sample. With dataset size  $m$ , the MSE is [33]:

$$J(\theta) = \frac{\sum_{i=1}^m (H_0(x^{(i)}) - y^{(i)})^2}{2m} \quad (4.4)$$



Seeing this equation already makes it clear that the number of training samples is important. From statistics, it is known that the population of data points, in this case training samples, have to be large enough. This problem is addressed later in Section 4.4.

### 4.2.3 Gradient descent

To solve the minimization problem, the first step is to start with  $\theta$  and keep changing the values to minimize  $J(\theta)$ . This is an iterative approach. Gradient descent is such an algorithm that can be used to find a solution to the minimization problem. Gradient descent is an algorithm that uses the gradient or derivative of a function to find a local minimum of that function.

In order to easily explain the gradient descent algorithm, an analogy to a ball rolling down a hill can be made. [34] Let's say we have a landscape such as in Figure 4.2. The horizontal location of the ball represents the values of  $\theta$ , the height of the ball represents the MSE. We want to get the ball to the lowest possible point, in order to minimize the MSE. To do this, the ball rolls down the hill until it comes to rest at the lowest point. Every iteration of gradient descent will bring the ball closer to the lowest point.

Of course, when a ball rolls down a hill, it does not just roll down and immediately stop. The ball takes a path similar as described in Figure 4.2. When the ball finally comes to rest, it can be said that the algorithm converges to that horizontal position.




FIGURE 4.2: Iterations of gradient descent. [35]

A gradient descent algorithm does this using the following algorithm with a simultaneous update for all values of  $\theta$  [33]:

repeat until convergence {

$$\theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \quad (4.5)$$

}

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\sum_{i=1}^m (H_0(x^{(i)}) - y^{(i)}) * x_j^{(i)}}{m} \quad (4.6)$$

The equation uses the derivative. A derivative indicates the rate of change, in this case, how steep the hill is. If the hill is very steep, then we move the ball a lot more than when the ground beneath the ball is almost flat.

$\alpha$  is the learning rate of the gradient descent. The value of  $\alpha$  describes how fast the gradient descent algorithm approaches the local minimum. If  $\alpha$  is too small, the gradient descent can be very slow. In the other case, if  $\alpha$  is too large, the gradient descent can overshoot the local minimum. It may fail to converge and could even diverge.

To go back to the example of the rolling ball, having  $\alpha$  too large, could be seen as having the ball roll down the hill extremely fast. Because of this, the ball could come to rest at another point in the landscape, it might have flown over the hill.

The value of  $\alpha$  does not need to change during the gradient descent, since the closer the gradient descent gets to the local minimum, the smaller the derivative becomes, and smaller steps will be taken. If  $\theta_j$  is already a local minimum, the derivative is 0 and the gradient descent will not change the value of  $\theta_j$ . [36]

#### 4.2.4 Normal equation

There is another method that could be used in place of gradient descent, a normal equation. This is a method to solve for  $\theta$  analytically. But this method becomes slow when there are a lot of features.  $X$  is a  $m \times n$  matrix constructed by putting all training samples (vectors of features) together. The notation  $x^{(i)}$  and  $y^{(i)}$  is used to denote the  $i$ th training sample. With dataset size is denoted by  $m$ .

$$\theta = (X^T X)^{-1} X^T y \quad (4.7)$$

But what happens when  $(X^T X)$  is non-invertible, or singular? This means there are redundant features or more features than training samples. There are mathematical models, such as the pseudo-inverse to still compute a correct result. There are still other methods that could replace gradient descent, such as conjugate gradient, BFGS and L-BFGS. Thinking about matrices and

non-invertibility can help to be able to understand how features relate to each other.

#### 4.2.5 Feature scaling

The main idea behind feature scaling is to make sure that the different features are on a similar scale. The reason behind this is to optimize the gradient descent algorithm. The gradient descent will converge much more quickly with feature scaling. With feature scaling, the partial derivative will be larger and because the features are in scale, the values of  $\theta$  will also be similar in scale.

When one feature is very large and another is very small, the values of  $\theta$  need to be small for the first feature and larger for the second, since the cost function needs to be minimized. This means that the values of  $\theta$  need to change more and thus, requires more iterations for the gradient descent algorithm. The range that should approximately be used is  $-1 < x < 1$ .

Mean normalization could be used. This replaces each  $x_i$  with  $\frac{x_i - \mu_i}{s_i}$ , where  $\mu_i$  is the average value of all  $x_i$  values and  $s_i$  is the standard deviation.

Lets use a numeric example. There are two data points,  $(0.5, 3, 1)$  and  $(0.7, 4, 2)$  with  $(x_1, x_2, y)$ . The hypothesis is:

$$H_0(x) = \theta_1 * x_1 + \theta_2 * x_2 \quad (4.8)$$

In Figure 4.3, it can be seen that the values of  $\theta$  quite quickly fall into the blue zone, where the cost function is minimal. The values of  $\theta$  also seem to be equally large when the cost function is minimal. This in contrary to Figure 4.4. Here, both values of  $\theta$  need to change a lot. They are not even remotely close to each other.




FIGURE 4.3: Cost function with feature scaling

When  $\theta_2$  is near  $-4$ , the gradient descent will change  $\theta_1$  towards a larger value. However, when  $\theta_2$  is near  $4$ , the gradient descent will change  $\theta_1$  towards a smaller value. In the gradient descent algorithm,  $\theta_2$  will converge towards approximately  $-1$  and the actual value of  $\theta_2$  will change between being larger and smaller than this value. Because of this, the value of  $\theta_1$  takes a longer time to converge.




FIGURE 4.4: Cost function without feature scaling

### 4.3 Classification with logistic regression

In a classification model, the machine learning algorithm tries to sort data into different classes. The simplest version is binary classification. For example, "is the email spam or not". In linear regression, the hypothesis can output values other than the classes that exist. However there is a method, logistic regression, which constrains the hypothesis to the available classes.




FIGURE 4.5: Binary classification. [37]

Figure 4.5 shows a simple binary classification example.  $x_1$  and  $x_2$  are features. Some points belong to the minus class, others belong to the plus

class. A logistic regression or classification algorithm will learn from training data what data from the minus class and data from the plus class looks like. Afterwards, the classifier can be used to predict whether an input sample belongs to the minus class or the plus class.

### 4.3.1 Logistic regression

In logistic regression, a hypothesis is needed which outputs whether an input datapoint belongs to a class or not. It can also be explained as, "what is the chance that the input belongs to a class"? Now the hypothesis needs to output a percentage, the chance that the input belongs to a certain class.

Then a decision boundary is used. This decides whether the input belongs to the class or not. Usually the decision boundary is 0.5. If the hypothesis outputs a value higher than the decision boundary, it can be said that the input belongs to the class for which it is being tested. Logistic regression uses the Sigmoid or logistic function for the hypothesis:

$$g(z) = \frac{1}{1 + e^{-z}} \quad (4.9)$$




FIGURE 4.6: Sigmoid function

Using this function, the hypothesis can be written as:

$$H_0(x) = g(\theta^T X) = \frac{1}{1 + e^{\theta^T X}} \quad (4.10)$$

### 4.3.2 Cost function

The cost function from linear regression cannot simply be applied to logistic regression. For classification it is important that data that belongs to the class has an output closer to 1, and that data that does not belong to the class has an output closer to 0. To do this a different cost function should be used. The notation  $x^{(i)}$  and  $y^{(i)}$  is used to denote the  $i$ th training sample. With dataset size is denoted by  $m$ . A helper function *Cost* is used. This helps to see the parallel with the cost function for linear regression, as seen

in Section 4.2.2. [38]

$$J(\theta) = \frac{\sum_{i=1}^m (Cost(H_{\theta}(x^{(i)}), y^{(i)}))}{m} \quad (4.11)$$

$$Cost(H_{\theta}(x^{(i)}), y^{(i)}) = \begin{cases} -\log(H_{\theta}(x^{(i)})), & \text{if } y^{(i)} = 1 \\ -\log(1 - H_{\theta}(x^{(i)})), & \text{if } y^{(i)} = 0 \end{cases} \quad (4.12)$$

The given cost function accomplishes exactly what is wanted. If the expected output is 1 (belongs to a class), then the cost becomes greater as the output goes to 0 and analogous if the expected output is 0. Gradient descent for logistic regression is exactly the same as for linear regression except with a different hypothesis.

A very simple example of the cost function would be the following. Take one data point, (3) with outputs 1. The hypothesis seen in Figure 4.7 and is:

$$H_0(x) = \frac{1}{1 + \exp \theta_1 * x_1} \quad (4.13)$$

The cost function is:  $-\log(\theta_1 * 3)$  which is graphed in Figure 4.8 becomes larger when the hypothesis goes towards 0, and goes to 0 when the hypothesis goes towards 1. This is indeed the behaviour that is expected from the cost function for logistic regression.




FIGURE 4.7: Hypothesis function example




FIGURE 4.8: Cost function example

### 4.3.3 Multi-class classification

The above hypothesis and cost function are for binary classification. To compute multi-class classification, the One-vs-all algorithm could be used. This algorithm splits the multiples classes into two groups. One group contains one class, the other group contains all other classes. Using these two new groups, the algorithms for binary classification can be used. In other words, for each class  $i$  a different logistic regression classifier  $H_{\theta}(x)$  is trained to predict the probability that  $y = 1$ . On an input  $x$ , the most probable class is chosen.

Another method called One-vs-One could be used. This method compares all pairs of classes. It checks, "does the input belong rather to class A as compared to class B". It does this for all pairs and decides which class most likely contains the input.

The One-vs-One method uses  $n_{classes} * (n_{classes} - 1)/2$  different logistic regression classifier  $H_{\theta}(x)$ . Each of these classifiers is trained to fit a pair of classes. [39]

To provide an example, Figure 4.9 can be used. On the left, binary classification is shown. The hypothesis and cost function as seen in the above sections can be used to do the classification. On the right, multi-class classification is used. There are three classes: triangles, squares and crosses.




FIGURE 4.9: On the left, binary classification is shown. On the right, three different classes are used. [40]

One-vs-all classification will have to use three classifiers to predict which class a given input sample belongs to. The first classifier predicts whether the input samples belongs to the triangles or to another class. The second classifier predicts whether the input samples belongs to the squares or to another class and analogous for the third classifier. This method can quite quickly find the result. If the first classifier is used first and it is found that the input sample belongs to the triangles class, the prediction is done. [40]

One-vs-One classification will have to use six classifiers. Each classifier

is trained to predict whether an input sample belongs to the triangles or the squares, to the triangles or the crosses, and so on. If the classifiers predict that the input samples belongs rather to the triangle class than to the square class and rather to the triangle class than to the crosses class, it can be predicted that the input sample belongs to the triangle class. Because the One-vs-One classification has to use more classifiers than One-vs-All, it is also slower. However, because it compares all pairs it is also more accurate. [40]

## 4.4 Overfitting

When the data is not modeled correctly and the model is too precise for the training data, a problem called overfitting occurs. Models that are overfitted have high variance, there are too many possible hypotheses. In other words, if there are too many features, the hypothesis may fit the training set very well but fails too correctly predict new examples. The hypothesis becomes very complex. In a similar way, underfitting may occur.




FIGURE 4.10: The points are generated by a sin function with Gaussian noise. Left: good fit (polynomial of degree 3), Right: overfit (polynomial of degree 10). [41]

There are methods to avoid overfitting. It is possible to reduce the amount of features. This can be done manually or by using a model selection algorithm which is explained in Section 7.2.3. Another option is regularisation. This method keeps all features but manages the values of the parameters of  $\theta$ . Regularisation works well when there are a lot of features that are all contribute to be able to predict  $y$ .

Regularisation is a method to resolve overfitting. When a hypothesis overfits, the hypothesis is too complex. This means that the hypothesis fits the training data very well and does not generalise well as seen in Figure 4.10. It fits the training data too good and as such the hypothesis cannot properly make predictions on other data.

A hypothesis is complex when the values of  $\theta$  are too large. Regularisation tries to solve this problem. When regularisation is used, the chosen



values of  $\theta$  become lower. For linear regression, this can be done in the cost function by redefining the cost function as:

$$J(\theta) = \frac{\sum_{i=1}^m (H_0(x^{(i)}) - y^{(i)})^2 + \lambda * \sum_{j=1}^m (\theta_j^2)}{2m} \quad (4.14)$$

The only change to the cost function is that a penalty is added for large values of  $\theta$ . When values of  $\theta$  become larger, the cost function also becomes larger. When this cost function is minimized lower values for  $\theta$  will be chosen. Only  $\theta_1$  till  $\theta_m$  should be regularised.  $\lambda$  is called the regularisation parameter. Because of this small modification, the hypothesis that is constructed will be simpler.

Lets use a numeric example. There are two data points, (3, 1) and (5, 2). The hypothesis is:

$$H_0(x) = \theta x_1 \quad (4.15)$$




FIGURE 4.11: Effect of regularisation

In this example, a cost function is constructed with and without regularisation. The blue graph shows the cost function without regularisation, the red graph shows the cost function with regularisation. It can be seen that having regularisation does change the value of  $\theta$  where the cost function is minimal, which is exactly what is intended to happen with regularisation.

## 4.5 Summary

At the core of a machine learning algorithm is the hypothesis. The hypothesis is the function that is used to fit the training data, it is used to predict values for new data. The hypothesis is constructed using the features that represent the data given to the machine learning algorithm and a set of unknown values  $\theta$ .

These values  $\theta$  are chosen, as such that the hypothesis is the most accurate. In order to do this, a cost function is constructed. The cost function measures the difference between a known output and a predicted output, usually the MSE is used. The values for  $\theta$  can be found by minimizing the cost function. This is done by using the gradient descent algorithm.

## Chapter 5

# Algorithms

This chapter gives an introduction to different machine learning algorithms. It starts with two popular supervised learning algorithms, Support Vector Machines and K-Nearest Neighbors. Afterwards it explains how neural networks work. Finally some unsupervised techniques such as clustering and anomaly detection are discussed.

### 5.1 Support Vector Machines

Support vector machines or SVMs is a supervised learning algorithm which offers an alternative view on logistic regression. Support vector machines try to find a model which divides 2 classes exactly with the same amount of margin on either side as shown in Figure 5.1. Samples on the margin are called the support vectors.




FIGURE 5.1: Support vector machines with their support vectors. [42]

#### 5.1.1 Linear and non-linear classifiers

Normally Support Vector Machines, as well as general logistic regression work with linear classifiers. A classifier can be linear or non-linear. A linear classifier has a decision boundary which is a linear function. This can be visualised by visualising that there is a line drawn through the data. On one side of this line, the data belongs to one class. Data on the other side of this line belongs to another class. [43]

Sometimes problems cannot be solved with linear classifiers. When training data contains a lot of noise for example. Noise in data could mislead the learning method and increase the classification error. Figure 5.2 shows classification of webpages being Chinese or English web pages. This seems to be a linear problem and a linear classifier has divided the space in the two classes (dashed line and short dashes). However, there are three noise documents (marked with arrows). The noise is misleading the classifier. It increases the classification error. [44]

However, when a non-linear classifier is applied, the noise is not a problem anymore. The noise elements belong to their respective classes. This classifier performs better than the linear classifier. In this example, the noise elements are considered important and that has as effect that the non-linear classifier is not considered overfitted. [44]




FIGURE 5.2: Linear and non-linear classifiers. [44]

### 5.1.2 Kernels

In order to adapt Support Vector Machines to be able to fit non-linear classifiers, some adjustments need to be done. This can be done with kernels. A kernel is a similarity function. The function compares two inputs and computes their similarity. This method, rather than preprocessing the entire feature space, all features at once, is called the kernel trick. [45].

The kernel trick is often computationally cheaper than the explicit computation of "new" features which represent data and their similarity. Normally features are extracted from data and then fed into a machine learning algorithm. Kernels offer an alternative.

The kernel should be a function to compare input data. The kernel, along with labeled data is then used to construct features. In literature, when no special kernel is used, the algorithm is said to use a linear kernel. The basic type of kernel algorithms are called Gaussian (or RBF) kernels. The formula for Gaussian (or RBF) kernels is:

$$K(x, y) = \exp\left(\frac{-\|x - y\|^2}{2\sigma^2}\right) \quad (5.1)$$

The  $\|x - y\|^2$  is the squared Euclidean distance between two features. The features are  $x$  and  $y$ .  $\sigma$  is a free parameter which can be tweaked when using the kernel to provide optimal results for the specific problem that is being solved.

## 5.2 K-Nearest Neighbors

The K-Nearest Neighbors or KNN algorithm is a supervised learning algorithm which computes the classification by looking at the classes of the K-Nearest neighbors of the training data. The K-Nearest Neighbors is an Instance-based algorithm. [46] The chosen class is the class most common among its K-Nearest neighbors, this can be seen in Figure 5.3. K is typically a small number. When K is equal to 1, the assigned class is the same as the class of the closest sample.




FIGURE 5.3: K-Nearest Neighbors. [47]

When training the algorithm, the input data and classes are stored. There are multiple methods to compute the distance between data. Euclidean distance can be used for continuous data. For discrete variables another metric can be used, such as the Hamming distance. There are many different methods for computing the distance.

### 5.2.1 Distances

There are several different methods to compute distances, both for continuous features and discrete features. These distance metrics are general mathematical constructs and not specific to the KNN algorithms. However, in machine learning, distance metrics are only used with KNN or derivative algorithms. They are generally not used in other machine learning algorithms.

**The Euclidean distance** is the distance between two points in the Euclidean space. This can be seen in Figure 5.4. With features  $x = (x_1, x_2, \dots, x_n)$  and  $y = (y_1, y_2, \dots, y_n)$ , the Euclidean distance is defined as  $\sqrt{\sum_{i=1}^n \|x_i - y_i\|^2}$ . When features are continuous, the Euclidean distance is a good metric. However, discrete features cannot be handled by the Euclidean distance. [48]

For example, when a feature represents something such as "country of origin". Countries could be considered continuous. For example the distance between countries could be calculated using  $(latitude, longitude)$ . In this example, the feature represents the "country of birth" of a person. In that case it does not make sense to use a continuous feature. In this case, the "country of birth" is either equal or not equal. This also shows that a single feature such as a country can be represented as either continuous or discrete depending on the problem that is being solved.

Features are almost always represented in numeral methods. "Belgium" could be 1, "Netherlands" could be 2, "France" could be 3, and so on. The difference between "Belgium" and "Netherlands" should be the same as "Belgium" and "France".




FIGURE 5.4: Euclidean Distance. [49]

**The Hamming distance** is defined as the number of positions at which the corresponding symbols are different.[50] For example, "John" and "Jack" have a Hamming distance of 3. "Jack" and "Jace" have a Hamming distance of 1. This can be seen in Figure 5.5 Going back to the example of

countries. Each country can be represented as a number. Each number is seen as a different symbol. For example 10 is seen as a single number. With the countries, the Hamming distance can only be 0 (countries are the same) or 1 (countries are different).




FIGURE 5.5: Hamming Distance. [50]

**The Chebyshev distance** is a distance which is defined as the largest difference between the features. With features  $x = (x_1, x_2, \dots, x_n)$  and  $y = (y_1, y_2, \dots, y_n)$ , the Chebyshev distance is defined as  $\max_{i=1} (||x_i - y_i||)$ . It is also known as chessboard distance. [51] It is the amount of moves a king needs to move to another space on a chessboard as seen in Figure 5.6.




FIGURE 5.6: Chebyshev Distance. [52]

**The Manhattan distance** is similar to the Euclidean distance. With features  $x = (x_1, x_2, \dots, x_n)$  and  $y = (y_1, y_2, \dots, y_n)$ , the Manhattan distance is defined as  $\sum_{i=1}^n ||x_i - y_i||$ . The Manhattan distance is the sum of the difference between each feature. [53] The Manhattan distance in a two-dimensional space is the sum of the difference of the  $x$ -values and the difference of the

$y$ -values. An example can be seen in Figure 5.7.




FIGURE 5.7: Manhattan Distance. [53]

**The Canberra distance** is a weighted version of the Manhattan distance. With features  $x = (x_1, x_2, \dots, x_n)$  and  $y = (y_1, y_2, \dots, y_n)$ , the Canberra distance is defined as  $\sum_{i=1}^n \frac{|x_i - y_i|}{|x_i| + |y_i|}$ . The Canberra distance is the same as the Manhattan distance divided by the sum of the different values. [54]

**The Minkowski distance** is a generalisation of the Euclidean distance and the Manhattan distance. Below the formulas for the Euclidean distance and the Manhattan distance can be seen. However, they are written slightly different. These formulas are almost the same.

$$EuclideanDistance(x, y) = \left( \sum_{i=1}^n \|x_i - y_i\|^2 \right)^{\frac{1}{2}} \quad (5.2)$$

$$ManhattanDistance(x, y) = \left( \sum_{i=1}^n \|x_i - y_i\|^1 \right)^{\frac{1}{1}} \quad (5.3)$$

A generalisation can be written as:  $\left( \sum_{i=1}^n \|x_i - y_i\|^p \right)^{\frac{1}{p}}$ .  $p$  is a free variable which can be set. If  $p$  goes to infinity, the Chebyshev distance is obtained. The effect of  $p$  is shown in Figure 5.8. This shows that the Minkowski distance is very flexible and powerfull. [55]






FIGURE 5.8: Minkowski Distance. [56]

All these distances are different methods that can be used with the KNN or derivative algorithm. The performance of each of these distances on the accuracy depends on the problem that is being solved. Finding the correct distance metric is a matter of trial and error.

### 5.2.2 KNN drawbacks & solutions

A major drawback of the KNN algorithm is the weakness to skewed data. Skewed data is data that belongs to a class that is underrepresented in the complete data set. Since the class is chosen based on the most popular nearest class, these popular classes may dominate the prediction. One simple solution is to gather more data. More realistically, the problem can be overcome by taking the distance between the input data and the neighbors into account. If this does not work, it is possible to use learning algorithms to select the best data from the data set so that no class is underrepresented. These methods are far more advanced and have a higher computational cost.

In Figure 5.10, an example is shown with 3-class classification. In this example, the weights are uniform. The score that decides which class is going to be predicted is computed from a majority vote. The example from Figure 5.11 assigns weights proportional to the inverse of the distance between the data point which is being predicted and any  $k$  neighbors. The distance function used to select the class is similar to  $\frac{1}{\text{distance between } x \text{ and } y}$  with  $x$  the data point for which a class is going to be predicted and  $y$  a neighbor of  $x$ . This is the method to address skewed data.

If  $k$  would be set to 2, then the 2 closest neighbors are chosen. The class is then chosen based on the inverse distance. Assume that the 2 closest neighbors belong to classes  $A$  and  $B$ . The first neighbors is closer to the data point for which a class is going to be predicted. Then the class that is chosen for the data point is  $A$ , not  $B$ . In Figure 5.9, it can be seen that when the distance (red line) increases, the weight decreases (blue line). All

weights that belong to neighbors of the same class are added together, the class with the highest score is chosen.




FIGURE 5.9: Inverse distance function




FIGURE 5.10: K-Nearest Neighbors example with 3 classes with uniform weights. [57]




FIGURE 5.11: K-Nearest Neighbors example with 3 classes with inverse distance weights. [57]

### 5.2.3 Selection of parameter

The best choice of  $k$  depends on the problem that is being solved and on the data that is being used to learn from. Larger values for  $k$  reduce the effect of noise data during prediction. However, the other effect of large values of  $k$  is that classes will have less distinct boundaries between them.

In general, the presence of noise severely limits the accuracy of the algorithm. Another issue that can limit the accuracy of the algorithm is irrelevant features. There are different methods to select and scale features in order to improve accuracy. One approach is to use evolutionary algorithms. [58]

## 5.3 Clustering

Clustering is a machine learning concept using unsupervised learning. Unsupervised learning does not have labels with the training set. An unsupervised machine learning algorithm tries to find structure within the given training set. Clustering is the first type of unsupervised learning. It tries to cluster the training samples into different clusters.

### 5.3.1 K-means Algorithm

The K-means algorithm is a simple clustering algorithm.  $K$  is the number of clusters that is going to be used. An example of how K-means could cluster data is shown in Figure 5.12. The algorithm first randomly places the  $K$  clusters in the space (from the training set). This can be done by randomly choosing  $K$  training samples. Then it repeats the following steps until the cluster centers remain stationary. it iterates over all training samples, and assigns them to the closest cluster. Next the cluster centers are moved to the center of the total cluster. [59]

---

```

# choose k random samples
c = choose_random_samples(k)
do:
    for sample in all_samples:
        c_i = closest_cluster_to(sample)
        c_i.append(sample)
    c.recalculate_center()
while cluster_center_changed(c)

```

---

The final clusters that are found by K-means are dependent on the random placement of the K clusters in the beginning. It could lead to suboptimal clustering. This can be fixed by running K-means a number of times and after each iteration check the value of the cost function to find the most efficient run of K-means. The cost function can be described as: [60]

$$J(c, \mu) = \frac{1}{m} \sum_{i=1}^m (\|x^{(i)} - \mu^{(i)}\|)^2 \quad (5.4)$$

There is a similar problem with determining the amount of clusters. The same method could be used to determine the correct number of clusters. However, manually determining the number of clusters could be more time efficient. [59]




FIGURE 5.12: K-means clustering. [61]

## 5.4 One-class Support Vector Machine

Support Vector Machines can also be used for unsupervised learning. The SVM will only have one class. The SVM can be trained using a dataset filled with data that belongs to that one class. The SVM learns a decision boundary around these data points as seen in the example from Figure 5.13. This boundary can be used to predict whether a data point belongs to that class or whether it does not. One-class SVM's can also make use of kernels.




FIGURE 5.13: One-class SVM. [62]

## 5.5 Neural networks

Neural networks are a useful alternative to logistic regression if the amount of features becomes too large. The origin of neural networks are algorithms which try to mimic the brain. There is a hypothesis, the "one learning algorithm" hypothesis, that shows that the brain can learn very different things, such as sound, touch, etc. by using a single algorithm. [63]

A neural network is created of neurons, which are called a logistic unit. Each neuron receives input wires, and has an output wire, which computes a value using the sigmoid (logistic) hypothesis, the activation function. A neural network is a group of "neurons" that are connected together. This can be grouped into a layered approach as seen in Figure 5.14. The first layer is called the input layer, the final layer is called the output layer which outputs a  $H_{\theta}(x)$  which is class. All layers inbetween these layers are called hidden layers. [64]




FIGURE 5.14: A neural network showing the different layers.[65]

Each logistic unit is denoted by  $a_i^j$ .  $j$  is the layer and  $i$  is the position in that layer.  $\theta^j$  is a matrix of weights or parameters controlling function mapping from layer  $j$  to layer  $j + 1$ .  $s_l$  is the number of units within a layer. The number of layers is denoted by  $L$ .

The neural network works similar to logistic regression, except that it performs logistic regression from layer to layer. The parameters  $\theta_j$  required are learned by itself. The architecture of a neural network refers to the way the units are connected to each other. Neural networks can be used for multi-class classification. Here, there are multiple units in the output layer and each unit represents a different class.  $K$  will denote the amount of output units. [66]




FIGURE 5.15: A neural network capable of multi-class classification. [67]

Data flows using the principle of forward propagation. The data passes through the first layers, moves to the second layers and so on, until it arrives at the final layer. Mathematically this can be described as:

$$\begin{aligned} a^{(1)} &= x \\ a^{(2)} &= g(\theta^{(1)} * a^{(1)}) \end{aligned}$$

$a^{(1)}$  is the input layer and  $x$  are the features of a data point that are fed to the neural network.  $g(\theta^{(1)} * a^{(1)})$  is the sigmoid logistic regression function with as input  $\theta^{(1)} * a^{(1)}$ .  $\theta^{(1)} * a^{(1)}$  can be defined as  $z^{(1)}$  to make the formulas simpler. This is a standard example of logistic regression. This happens for every node in a layer and for every layer. Every layer takes as input, the output from the previous layer. The data propagates forward through the network. [68]

### 5.5.1 Cost function and backpropagation

The cost function for neural networks is a generalisation of the cost function for logistic regression. The cost function accounts for the different layers, units and the number of output units. To minimize the cost function, the same methods such as gradient descent can be used. However, the problem is how to compute the partial derivative of the cost function. Using backpropagation, it is possible to compute this. [69]

The resulting class is known in the final layer. From here, the algorithm can find the error.  $\delta_j^l$  will be the symbol used for the error of node  $j$  in layer  $l$ . For an example with 4 layers:

$$\begin{aligned} \delta_j^{(4)} &= a_j^{(4)} - y_j \\ \delta^{(3)} &= (\theta^{(3)})^T \delta^{(4)} * g'(z^{(3)}) \end{aligned} \tag{5.5} \tag{5.6}$$

The algorithm starts by setting a parameter  $\Delta_{ij}^{(l)}$  to 0 for all  $i, j, l$ . Then, it iterates from  $i = 0$  to  $m$ , with  $m$  the number of training samples. Each iteration,  $a^{(1)}$  is set to  $x^{(i)}$ . Forward propagation is computed for  $a^{(l)}$  for all  $l = 2, 3, \dots, L$ . Using  $y^{(i)}$ ,  $\delta^{(L)}$  can be computed. Then the different  $\delta^{(L-1)}$  to  $\delta^{(2)}$  are computed. Finally,  $\Delta_{ij}^{(l)}$  is incremented by  $a_j^{(l)} \delta^{(l+2)}$  for each  $l$ . After the iteration is done, the final value for the partial derivative can be calculated:  $\frac{\partial}{\partial \theta_{ij}^{(l)}} J(\theta)$ . This can be done by dividing  $\Delta_{ij}^{(l)}$  by the amount of training samples and adding  $\lambda \theta_{ij}^{(i)}$ . [70]

### 5.5.2 Using a neural network

A neural network should have as many input units as the dimension of features. The number of output units is equal to the number of classes. Default, there should be either 1 hidden layer or if there are more, all layers should have the same number of hidden units. [71] The neural network should be trained by first assigning random weights to the values of  $\theta$ . Afterwards, forward propagation should be used to get  $H_\theta(x^{(i)})$ . Next the cost function should be computed. Backwards propagation is used to compute the

partial derivatives. The result of backwards propagation can be checked by numerical methods to compute the gradient. Finally gradient descent or other advanced optimization methods with backpropagation should be used to try to minimize the cost function as a function to the parameters  $\theta$ .

A neural network could be used to implement Autonomous Driving. The neural network can be trained using images which display a road. [72] The labeled data could be classes explaining whether the road goes straight, left or right. This could be used to make vehicles drive on a road, but could also be used to program drones to be able to navigate through forest paths. [73]

### 5.5.3 Deep neural networks

Deep Learning Algorithms are a modern modification to Artificial Neural Networks that exploit abundant cheap computation. Deep learning networks are very deep and complex neural networks. They contain a lot of hidden layers. Deep Boltzmann Machine (DBM), Deep Belief Networks (DBN) and Convolutional Neural Network (CNN) are examples of deep learning algorithms. These algorithms have specific rules about how the network is connected. [74]

## 5.6 Decision tree algorithms

Decision tree algorithms are a set of supervised learning algorithms used for classification and regression. The algorithms try to create a model that resembles a tree. Every prediction starts at the root of the tree. At every node, a certain decision rule is used. A decision rule is a condition, for example it could be: "Is feature A larger than 10". Once a leaf of the tree is reached the correct class or value is found. [75] An example of this can be seen in Figure 5.16



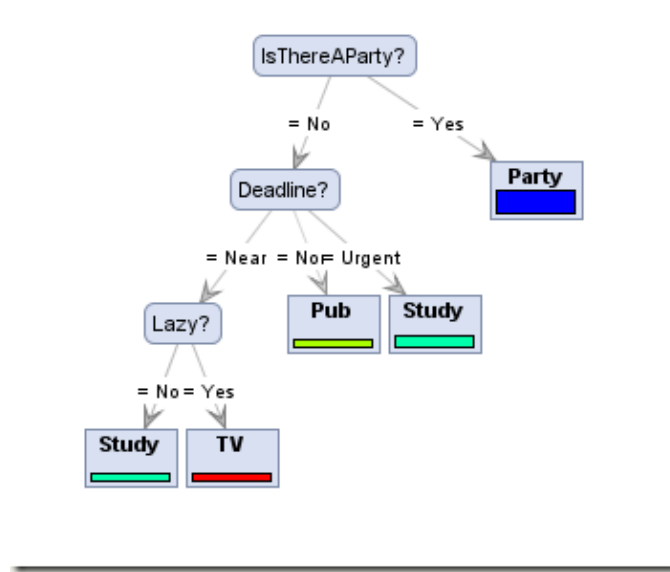


FIGURE 5.16: Decision Tree Algorithm example [76]

A decision tree is simple to understand and to visualise. They are also quite fast. However, a decision tree algorithm can easily create complex trees and overfit. They also cannot handle skewed data that well. [75]

## 5.7 Bayesian Algorithms

There are Bayesian Algorithms which explicitly apply Bayes' Theorem for problems such as classification and regression. There are a lot of variances on Bayesian Algorithms, the most popular is Naive Bayes. [46]

Bayes Theorem is stated as follows [77]:

$$P(A|B) = \frac{P(B|A) * P(A)}{P(B)} \quad (5.7)$$

$A$  and  $B$  are events.  $P(A)$  is the chance that event  $A$  happens and  $P(B)$  is the chance that event  $B$  happens.  $P(A|B)$ , a conditional probability, is the probability of event  $A$  given that  $B$  is true.  $P(B|A)$  is the probability of event  $B$  given that  $A$  is true. [77]

Naive Bayes is a supervised learning algorithm with the "naive" assumption that every feature is independent. The formula for the Bayes' Theorem is used for Naive Bayes. The event  $A$  is a class.  $B$  is the collection of features. Naive Bayes is decent for classification but is not optimal for regression. [78]

## 5.8 Association Rule Learning Algorithms

Association Rule Learning Algorithms are methods that extract rules that best explain observed relationships between variables in data. Apriori algorithm and Eclat algorithm are examples of such algorithms. These algorithms operate mainly on databases. [46]

In Figure 5.17, an example is given of a small database. Association Rule Learning could make several rules from this database. A rule could be "if a set contains alpha and beta, then there is a 50% chance that it also contains epsilon". [79] It does this by counting items. First it counts the frequency that alpha, beta, etc. appears in the database. Next, it counts the all the pairs (*alpha*, *beta*). The algorithm continues until it cannot make any more sets. From these sets, it can make the association rules.

alpha	beta	epsilon
alpha	beta	theta
alpha	beta	epsilon
alpha	beta	theta

FIGURE 5.17: Apriori Algorithm example [79]

## 5.9 Specific algorithms for a sub-field

A lot of algorithms are specifically constructed for a specific sub-field of machine learning, for example computer vision, natural language processing, etc. Even within the categories of algorithms that were discussed, regression, regularization, instance-based, clustering and neural networks, there are a lot of different variants. However, these are considered to be too advanced and outside of the scope of this thesis.

## Chapter 6

# Machine Learning Techniques

### 6.1 Dimensionality reduction

Dimensionality reduction is the process of reducing the amount of features used in machine learning algorithms. This can be used to increase the accuracy and the performance of machine learning algorithms. One form is to do data compression. For example, transform 3D data into 2D data and eliminate a feature or dimension. It can also be used to reduce dimensions to be able to efficiently visualise data.

Dimensionality reduction can be used to speed up the time it takes for other learning algorithms to learn. By using dimensionality reduction, the amount of features or the amount of training samples is reduced which reduces the running time of the training, but the compressed data still retains the same information as the uncompressed data.

#### 6.1.1 Principle Component Analysis

Principle Component Analysis is a way to do dimensionality reduction. The algorithm is formulated as a minimisation problem. When given  $N$ -dimensional data and  $N-1$  dimensional data is preferred, the algorithm tries to find the correct  $N-1$  dimensional value so that the projection is the closest to the original data.

Before this algorithm is run, the features of the data should be scaled, so that all features are on a similar scale. This can be done by using mean normalization. In order to reduce the dimension from  $n$  to  $k$  the covariance matrix should be computed.

From this matrix, the eigenvectors need to be computed using singular value decomposition. From these values, only the first  $k$  values are going to be used and be multiplied with the training data.

### 6.2 Anomaly detection

Anomaly detection is also a form of unsupervised learning. The algorithm learns what normal behaviour looks like through the training set and then tries to predict if a given input data belongs to the normal behaviour or is abnormal for any reason.

In Figure 6.1, the algorithm has been trained using the green data points.

These make up the normal behaviour. The algorithm can then predict whether a data point is normal, or an anomaly. It seems very close to One-class SVM's as seen in Section 5.4. However, One-class SVM's are used for simple binary classification. Normal behaviour of a systems typically does not fall within a single class. Usually many different classes make up normal behaviour.




FIGURE 6.1: Anomaly detection. [80]

Anomaly detection is useful when there are a lot of data points belonging to normal behaviour and there are almost no abnormal behaviour data points. General supervised learning algorithms are usefull when there are a lot of data points of both normal and abnormal behaviour. [81]

### 6.2.1 Normal Distribution

Anomaly detection algorithms make heavy use of (Gaussian) Normal distribution:

$$x \sim N(\mu, \sigma^2) \quad (6.1)$$

Hereby  $\mu$  is the mean parameter and  $\sigma$  is the standard deviation. Now the probability of  $x$  being an anomaly can be calculated as:

$$p(x) = \prod_{j=1}^n (p(x_j; \mu_j, \sigma_j)) \quad (6.2)$$

$$p(x_j; \mu_j, \sigma_j) = \frac{1}{\sqrt{2\pi}\sigma_j} \exp\left(-\frac{(x_j - \mu_j)^2}{2\sigma_j^2}\right) \quad (6.3)$$

$p(x_j; \mu_j, \sigma_j)$  is the probability of  $x_j$  with a Normal Distribution with parameters  $(\mu_j, \sigma_j)$ .  $p(x)$  is the product of all these probabilities. The formula for

the probability is the probability density of the Normal distribution. This function is plotted in Figure 6.2. The  $y$ -axis is the density function, the  $x$ -axis is the  $x$  value in the formula.




FIGURE 6.2: Normal Distribution. [82]

### 6.3 Online learning

Online learning is a technique to update an algorithm. Instead of feeding all data to the learning algorithm, the data is fed incrementally. The algorithm can keep learning from data even after it has initially learned a model.

Some algorithms such as KNN are immediately able to do online learning. Other algorithms, such as SVM need to be slightly altered in order to be able to do online learning.

Online learning is very useful when the data is dependent on the time. For example, stock prices prediction. The data that is used to train an algorithm today might be able to predict stocks for tomorrow, but it is not very effective at predicting the stock market at long term. That is why it is useful to be able to keep training the algorithm. [83]

### 6.4 Bagging

Bagging is a technique which tries to create new datasets from a given data set. It takes a data set as input and generates multiple slight variations of the data in the given data set. These slight variations are used to construct new data sets. This method is also called Bootstrapped Aggregation, or bagging for short.

Once multiple data sets have been generated, it is possible to combine multiple weaker models and training algorithms and try to combine the predictions made by these algorithms. This is useful to deal with unstable data.

Unstable data is data that can give different results depending on the algorithm that is used to process the data. An example of bagging and Neural Networks can be seen in Figure 6.3. [46]




FIGURE 6.3: Bagging example with Neural Networks. [84]

## Chapter 7

# Validating an algorithm

### 7.1 Machine learning evaluation

Machine learning diagnostics are tests that can be run to get to know what is and isn't working with a machine learning algorithm. They also provide guidance as to how performance could be improved.

### 7.2 Error analysis

When trying to use machine learning for any purpose, such as intrusion detection systems, there are several considerations to be made. Another important step is the approach used to find the correct algorithm. The recommended approach is to start with a simple algorithm and test it with cross validation data. Afterwards, learning curves could be plotted to decide if more or less data, more or less features, ... are likely to help. Finally, error analysis can be done by manually examining the samples on which the algorithm made mistakes. This could help to spot any systematic trends in the type of samples on which the algorithm is making mistakes.

The error analysis mostly consists of manual or brute force work. The samples on which the algorithm is wrong need to be categorized based on which features could help it categorize correctly and on which class it belongs to. Calculating statistics, such as the accuracy of the algorithm can also help with error analysis. It is possible to brute force this approach by automatically trying a lot of different combinations of features. This is less efficient than manually tweaking the algorithm.

Another issue to account for are skewed classes. Skewed classes are classes that are underrepresented in training data. For example, in binary classification, when trying to classify a flow as malicious or not, the training set might only provide 0.5% malicious data. Knowing this, having an accuracy of 99% does not seem that great.

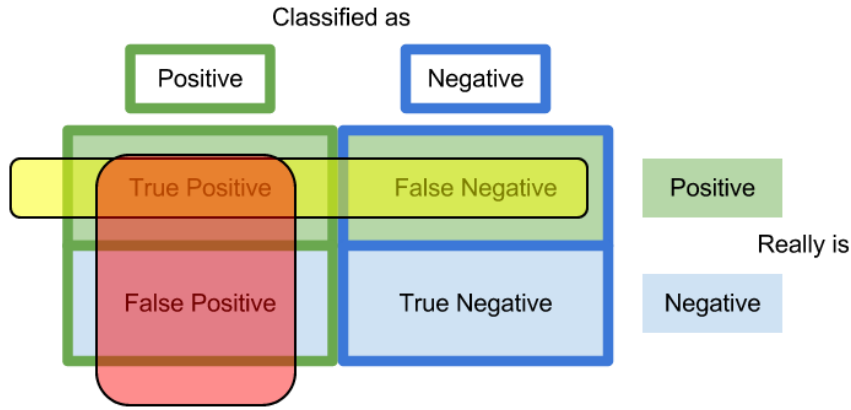


FIGURE 7.1: Precision and recall. [87]

A different metric is required to evaluate machine learning algorithms that are trained using skewed data. This metric is the precision and recall method. We classify a result as true positive, true negative, false positive and false negative. False positive and false negative respectively mean that the predicted value is falsely classified as positive and negative. These are the errors. True positive and true negative are correctly predicted values. **Precision** is defined as the fraction of predicted malicious flows that were actually malicious:

$$\frac{\text{truepositive}}{\text{truepositive} + \text{falsepositive}}$$

**Recall** is defined as the fraction of predicted positives and the actual amount of positives:

$$\frac{\text{truepositive}}{\text{truepositive} + \text{falsenegative}}$$

There is always a tradeoff to be made between precision and recall. As an effect of a higher precision, there will be a lower recall. Similarly, a higher recall means a lower precision.

Algorithms with a different precision and recall can be compared to each other. This can be done using an F-score:

$$2 \frac{PR}{P + R}$$

The F-score is the harmonic mean of precision and recall. For Table 7.1, this means that Algorithm 1 is the most effective. In contrast, using for example the average of both precision and recall would make Algorithm 3 the most effective.

The main reason to use the harmonic mean is because the average is taken of ratios (percentages), and in that case the harmonic mean is more appropriate than the (arithmetic) mean. [88]



TABLE 7.1: Example of precision and recall of certain algorithms.

	Precision (P)	Recall (R)
Algorithm 1	0.5	0.4
Algorithm 2	0.7	0.1
Algorithm 3	0.02	1.0

Looking at the values of the precision and the recall in Table 7.1, it can already intuitively be seen that Algorithm 1 is the best algorithm. For algorithm 3, the precision is 0.02, which means that there were a lot more false positives than true positives. With a recall of 1.0, there are almost no false negatives. This means that almost every prediction was a positive. That is not good at all.

Algorithm 2 has similar problem but the other way around. Only Algorithm 1 seems to have some kind of balance between precision and recall which makes it seem as the best algorithm. This confirms that the F-score is a better metric than the average.

TABLE 7.2: Example of average and F-score of certain algorithms.

	Average	F-score
Algorithm 1	0.45	0.444
Algorithm 2	0.4	0.175
Algorithm 3	0.51	0.039

The data that is being fed into a machine learning algorithm is also important to consider. Sometimes a lot of data can be useful. First, the assumption is made that the features are chosen correctly and sufficiently. Lots of data is useful when a machine learning algorithm is used with a lot of parameters such as a neural network with a lot of hidden units. This means that the algorithm has low bias.

### 7.2.1 Different Algorithms

If the number of features is large relative to the number of training samples, then using a linear kernel Support Vector Machine or logistic regression is preferred. A Gaussian kernel is preferred when there are more training samples than features, but the difference is rather small. Otherwise, new features should be added. Neural networks are almost always effective but they may be slower to train.

### 7.2.2 Evaluating the hypothesis

The most obvious way to test whether the hypothesis is correct is by dividing the training set into two sets. The first set which should have approximately 70% of the samples of the original training set will be used to train the learning algorithm. The other set is used to check whether the output of

the learning algorithm is correct. After the algorithm has done its learning with the training set, this check set is thrown into the algorithm. Afterwards the output of the algorithm is compared to the labels of the check set. If the output isn't correct, a test error can be computed. These test errors can be used to compute global error value, which could be calculated by, for example, a mean squared error. This value can be used to evaluate the hypothesis. [85]

### 7.2.3 Model selection algorithm

To get back to the problem of overfitting, there is another method next to regularisation called model selection. Model selection uses the same principle as mentioned above except it divides the training set into three new sets. A new training set, a cross validation set and a test set.

The training set is used to actually train the algorithm. The cross validation set is used to compare different algorithms or different models. Different models can be tested and compared to each other by comparing the cross validation error, the error between the prediction and the actual output from the cross validation set. The testing set is used purely to test the algorithm once the cross validation has been done. It is considered good practice to use separate sets for cross validation and testing. [86]

### 7.2.4 Diagnosing bias vs variance

High variance means that there is an overfitting problem. High bias on the other hand, is the opposite problem. It means that the hypothesis does not fit the training set at all. There is a relation between the amount of training samples and the training error. With more training data, the training error goes down.

With a low model complexity, there is a high error for the training set and a high error on the test set. This is a signal there is a underfitting problem or a high bias problem. It does not help to increase the amount of training samples since the model just is not complex enough.

With a high complexity model, the training set error is very low, but the test set error is very high. This is a signal there is an overfitting problem or a variance problem. This can be seen in Figure 7.2. The complexity of the model can be adjusted by changing the amount of features that are being used. [33]




FIGURE 7.2: The effect of the amount of training samples on the accuracy.

In Figure 7.2 the general trend of training error and test (or true) error is shown. This figure shows the example for high bias and no matter how many training samples there are, the functions are already converged to the same error amount.

In contrary to high bias, high variance can be solved by using extra training samples. In that case there is a gap inbetween the true error line and the training error line and it is likely that they still converge to the same point. With a bigger training set, they converge more and more. These curves are called learning curves.

## Chapter 8

# Machine learning for an IDS

An intrusion detection system has to detect whether some data it receives is either malicious or regular web traffic. This can be seen as a classification problem which means a machine learning algorithm for classification could be used. It needs to be determined whether data is either normal network traffic or, whether it corresponds to malicious behaviour.

Some parameters have to be chosen that will be feed into the machine learning algorithm. These parameters can be log files, entire packets, packet headers or IP flows as seen in Section 9. In general there are several advantages and disadvantages to using machine learning for intrusion detection.

### 8.1 Properties of using ML for an IDS

As said before, machine learning for an intrusion detection system is a classification problem. More precisely, it can be said that intrusion detection systems have to detect abnormal behaviour in a network with mostly normal behaviour. There are several problems that can be encountered when using machine learning techniques.

#### 8.1.1 Network diversity

A first problem is the diversity of network traffic. The notion of "normal network traffic" is difficult to actually define. The bandwidth, duration of connections, origin of IP addresses, applications used can vary enormously through time. This makes it quite difficult for machine learning algorithms to distinguish between "normal network traffic" and malicious behaviour.<sup>[89]</sup>

However, this issue can be bypassed by defining what malicious behaviour looks like. If a machine learning algorithm learns what malicious behaviour looks like, it does not need to know anything about normal behaviour.

#### 8.1.2 Detect new attacks

Another problem is the ability to detect new attacks. A machine learning algorithm compares incoming data with a model that it has created internally. An new type of malicious behaviour might appear to be closer to normal network traffic as compared to the model of known attacks. This is only an issue for completely new attacks since most attacks are similar to known

malicious behaviour. The model that is created for malicious behaviour can better detect newer attacks than a signature based IDS. [89]

### 8.1.3 Classification

There is only one big remaining problem when trying to teach a machine learning algorithm what malicious behaviour looks like. Different types of malicious behaviour can look completely different. When the machine learning algorithm constructs a model, the model might not be able to generalise well or it might give a lot of false positives or false negatives.

There are several solutions that can be used in order to make machine learning algorithms more effective for intrusion detection systems. One option is to change the way the classification problem is defined. Instead of defining the classes, "normal" and "malicious", there might be different classes for different types of malicious behaviour. In the same way, different classes can be defined for different types normal traffic.

This means that a different class is used for each different type of attack. This also makes it easy to define different severity levels for different attacks. Since the machine learning algorithm outputs the predicted class, a more detailed report can be generated on the anomaly. [89]

### 8.1.4 Learning

A big advantage of using machine learning is the learning part of the algorithm. The algorithm learns what malicious behaviour looks like. Whenever it has to predict whether something is malicious and what class it belongs to, it does not need to use perfect matching or heuristics such as other forms of intrusion detection as seen in Section 2.3. It has learned a model that describes what the different classes of malicious behaviour look like. [90]

## 8.2 Evaluating ML for an IDS

With a machine learning algorithm, performance can be measured using the F-score as seen in Section 7.2. However, for intrusion detection systems, this is not enough by itself. The F-score assumes that recall and precision has the same importance. This is not necessarily the case when evaluating intrusion detection systems. [90]

In Table 8.1, it is shown when there is a true positive and true negative. A false positive occurs when a sample is actually Normal but is classified as an Intrusion. A false negative occurs when a sample is actually an Intrusion but is classified as Normal.

TABLE 8.1: Performance measure.

	Intrusion	Normal
Intrusion	True Positive	False Negative
Normal	False Positive	True Negative

A false negative is bad, since it means that an Intrusion was not detected. But most IDS's are used in a layered approach. This means that if one layer does not detect an Intrusion, another layer might detect it.

The top layer usually looks at the entire network traffic. Because the entire network traffic might be huge, the data that is processed is kept to a minimum. This means that the top layer will use IP flows. This means that not detecting an Intrusion is not that horrible. A low recall could be disastrous if the amount of detected Intrusions is huge, which can happen if the amount of data that is passed through the IDS is also huge.

The layered approach might also work completely different. Perhaps the first layer tries to detect as many anomalies as possible (and having a low recall) and then passing the data for which anomalies have been detected to other layers. This approach means that a low recall is not bad.

The scoring used for an IDS that uses machine learning is dependant on how the IDS is going to be used. [90]

### 8.3 Using ML for an IDS

Data has to be processed before it can be used within a machine learning algorithm. This means that features have to be chosen. Some features can be easy to find, other have to be found by experimenting and running tests.

Using all the features of a dataset does not necessarily guarantee the best performances from the IDS. It might increase the computational cost as well as the error rate of the system. This is because some features are redundant or are not usefull for making a distinction between different classes. [91]

## Chapter 9

# IP Flows

Flows are aggregated from all packet data that travels through the network. Flow exporters are programs which collect network packets and aggregate them into flow records. A flow is not the same as a TCP connection. A flow can be any communication between two devices with any protocol. Flows are defined using a (source\_IP, destination\_IP, protocol) tuple. This is why flows are also called IP Flows.

Since flow data does not contain any payload information, intrusion detection systems that use flow data cannot detect malicious behaviour embedded within payload data. [13] A flow is constructed using a flow exporter. These constructed flows are usually collected and stored within a flow collector. An example of this can be seen in Figure 9.1.




FIGURE 9.1: Netflow. [92]

### 9.1 Attributes of a flow

The following attributes are available within a flow:

- Source IP
- Destination IP
- Protocol name

- Source port
- Destination port
- Starting time of the flow
- Duration of the flow
- Amount of packets in the flow
- Amount of bytes in the flow

These attributes are most commonly available within a flow. However there are different standards for a flow. Some contain less information, some contain more information. Should a flow exporter be implemented, some additional features can be generated from packet data.

- Amount of TCP SYN within the flow
- Source and Destination Type of Service
- TCP flags set

These extra attributes can be very usefull when trying to determine what is normal and abnormal behaviour.

The **source IP** is the IP address of the device or network that started the flow. The **destination IP** is the IP address of the device or network that the sender connected to. The IP address can be either an IPv4 or an IPv6 address. If the flows were gathered and created on a host computer, the IP address field can be a MAC address instead of an actual IP address. These type of flows are generated by newer flow exporters.

The **source port** is the port from which the flow originates. The **destination port** is the port to which the flow connects to.

The **protocol name** is the name of the protocol that is used for the flow. Each flow uses only one protocol. Some flow exporters might provide the protocol in numerical form, others in ASCII format. The numerical form is used mostly by flow exporters that only allow explicit IP traffic. Other flow exporters also contain non-IP traffic such as ICMP messages.

Flow attributes such as the **amount of packets** and the **amount of bytes** are self-explanatory. They are usefull to know for example whether the flow contains many packets which each contain almost no payload, or whether the flow contains one gigantic packet.

The **starting time** is the time when the flow is started. This is saved in unix time. The **duration of the flow** is the time between the start and end of a connection. When the flow is a TCP connection, the start and end times are defined. These are the moments that a SYN is sent and a FIN ACK is received. However, should the flow be an UDP flow for example, there is no clear end. The flow terminates itself after a timeout.



## 9.2 Using IP flows

In order to be able to use IP flows, data has to be extracted from the flows. In general, the different attributes from a flow are already useful. However it could also be interesting to extract other information from the flows.

### 9.2.1 Country of origin

The country of origin is something that can be extracted from flow information. The source and the destination IP can be given to an IP API which gives back information about the physical location of the IP address. Also the region, city name and ISP can be extracted. An example of such a query can be seen in Figure 9.2. This API is just one example of the available APIs on the web. [93]

```
{
  "status": "success",
  "country": "United States",
  "countryCode": "US",
  "region": "CA",
  "regionName": "California",
  "city": "San Francisco",
  "zip": "94105",
  "lat": "37.7898",
  "lon": "-122.3942",
  "timezone": "America/Los_Angeles",
  "isp": "Wikimedia Foundation",
  "org": "Wikimedia Foundation",
  "as": "AS14907 Wikimedia US network",
  "query": "208.80.152.201"
}
```

---

FIGURE 9.2: IP API JSON format. [93]

### 9.2.2 Time

The starting time can also be manipulated. From the unix starting time, information such as the year, month, day, weekday, hours and so on can be found. This information makes it easier to categorize network traffic and allows certain patterns to be seen. For example, it might be that certain types of web traffic only happen at a certain weekday.

## Chapter 10

# Implementation

### 10.1 Technology stack

The main library that was used is *scikit-learn* [94]. *Scikit-learn* is a machine learning library for Python. It is build upon NumPy, SciPy, and matplotlib. It is also open source and commercially usable with the BSD license.

Most algorithms that have been explained in Section 5 are implemented in *scikit-learn*. Neural networks, an promising algorithm, is not implemented in this library and has not been used in this thesis.

*Scikit-learn* also contains different methods to visualise machine learning algorithms such as a graph to show the learning curve. These can be a useful tool to evaluate the performance of machine learning algorithms.

#### 10.1.1 Program execution

The implementation works in different steps. A JSON config file is used to define the elements that are used within the program. This contains the data to be used for learning, for checking, the machine learning algorithm, etc. A full explanation of the config file as well as a general user guide can be found in Appendix A.

Once the config file has been read, the program can start the training phase. In this phase the specified algorithm is used and trained using the given data. Afterwards the prediction phase starts. This phase uses the prediction data and gathers all results.

#### 10.1.2 Structure

The implementation is build to be modular. New components can easily be added. Details on how to add new components to these modules is explained in the developer guide in Appendix B. The first module is the machine learning module. This module contains all machine learning algorithms that can be used.

There is also a feature module. This module contains the available classes that can be used to extract features from the flows. A loader module contains all classes required to load the data from the different datasets.

A training module contains the different classes used for training. These

classes use a loader class and pass the data to the machine learning algorithm. They define which data is supposed to be used (for example, using only abnormal behaviour and leaving out the normal behaviour).

Finally there is a results module. This module receives all the output from the machine learning algorithms and has to log these or visualise them.

## 10.2 Datasets

In order to test the implementation and the algorithms, different datasets were used. Each dataset is used to test a different aspect of the machine learning algorithms. First, a subset of a dataset has to be chosen to be fed to the machine learning algorithms for learning. Afterwards, using the method seen in Section 7.2.2, the algorithm is tested using another subset of the same dataset.

In the third step, the algorithms are tested using real-world data that is labeled. Finally, the algorithms are tested using raw, unlabeled real-world data. This is to make sure that the algorithm performs well on unprocessed real-world data. Several datasets have been used to test the machine learning algorithms.

### 10.2.1 CTU-13 Dataset

The CTU-13 dataset has been used for steps one to three for the testing of the machine learning algorithms. This is a labeled dataset. It contains botnet behaviour, normal and background traffic. The data was captured in the CTU University, Czech Republic, in 2011. It consists of thirteen different captures, each of which run a different botnet malware. Figure 10.1 shows the amount of data within each capture. Note that the captured data is only from a couple hours. [95]

Id	Duration(hrs)	# Packets	#NetFlows	Size	Bot	#Bots
1	6.15	71,971,482	2,824,637	52GB	Neris	1
2	4.21	71,851,300	1,808,123	60GB	Neris	1
3	66.85	167,730,395	4,710,639	121GB	Rbot	1
4	4.21	62,089,135	1,121,077	53GB	Rbot	1
5	11.63	4,481,167	129,833	37.6GB	Virut	1
6	2.18	38,764,357	558,920	30GB	Menti	1
7	0.38	7,467,139	114,078	5.8GB	Sogou	1
8	19.5	155,207,799	2,954,231	123GB	Murlo	1
9	5.18	115,415,321	2,753,885	94GB	Neris	10
10	4.75	90,389,782	1,309,792	73GB	Rbot	10
11	0.26	6,337,202	107,252	5.2GB	Rbot	3
12	1.21	13,212,268	325,472	8.3GB	NSIS.ay	3
13	16.36	50,888,256	1,925,150	34GB	Virut	1

FIGURE 10.1: Amount of data and botnet type for each capture. [95]

Each capture contains only a small amount of botnet samples as seen in Figure 10.2. Most flows are background flows. This is expected of botnet behaviour since it does not generate an large amount of network traffic. The actual labeling is much more detailed as compared to Figure 10.2. Each flow is labeled with its exact source. This could be google analytics, google web-mail or a windows update. The flows within the dataset only contain the regular information that is found within netflow. The abnormal behaviour within this dataset is internal abnormal behaviour.

Scen.	Total Flows	Botnet Flows	Normal Flows	C&C Flows	Background Flows
1	2,824,636	39,933(1.41%)	30,387(1.07%)	1,026(0.03%)	2,753,290(97.47%)
2	1,808,122	18,839(1.04%)	9,120(0.5%)	2,102(0.11%)	1,778,061(98.33%)
3	4,710,638	26,759(0.56%)	116,887(2.48%)	63(0.001%)	4,566,929(96.94%)
4	1,121,076	1,719(0.15%)	25,268(2.25%)	49(0.004%)	1,094,040(97.58%)
5	129,832	695(0.53%)	4,679(3.6%)	206(1.15%)	124,252(95.7%)
6	558,919	4,431(0.79%)	7,494(1.34%)	199(0.03%)	546,795(97.83%)
7	114,077	37(0.03%)	1,677(1.47%)	26(0.02%)	112,337(98.47%)
8	2,954,230	5,052(0.17%)	72,822(2.46%)	1,074(2.4%)	2,875,282(97.32%)
9	2,753,884	179,880(6.5%)	43,340(1.57%)	5,099(0.18%)	2,525,565(91.7%)
10	1,309,791	106,315(8.11%)	15,847(1.2%)	37(0.002%)	1,187,592(90.67%)
11	107,251	8,161(7.6%)	2,718(2.53%)	3(0.002%)	96,369(89.85%)
12	325,471	2,143(0.65%)	7,628(2.34%)	25(0.007%)	315,675(96.99%)
13	1,925,149	38,791(2.01%)	31,939(1.65%)	1,202(0.06%)	1,853,217(96.26%)

FIGURE 10.2: Distribution of labels in CTU 13 Dataset. [95]

### 10.2.2 Tracelabel Dataset

This dataset has been used for steps one to three for the testing of the machine learning algorithms. The tracelabel dataset is a database constructed at the University of Twente in September 2008. It contains flow data of traffic collected from a honeypot that was positioned in the network of the university. The honeypot ran several network services such as ftp, ssh, http, etc. The honeypot only captured abnormal behaviour. As such, most of the data within the dataset is actual abnormal behaviour. [96]

The flows within the dataset contain extra information. Each flow also contains the TCP flags obtained by OR-ing the TCP flags field of all packets of the flow. In Table 10.1, the labeling that is used within this dataset can be seen. The abnormal behaviour within this dataset is external abnormal behaviour.

TABLE 10.1: Labeling of the Tracelabel dataset.

Id	label	Explanation
1	ssh_scan	An ssh scan
2	ssh_conn	Unauthorized ssh connection attempt
3	ftp_scan	An ftp scan
4	ftp_conn	Unauthorized ftp connection attempt
5	http_scan	An http scan
6	http_conn	Unauthorized http connection attempt
7	authident_sideeffect	Unauthorized identification attempt
8	irc_sideeffect	Suspicious irc traffic
9	icmp_sideeffect	Suspicious icmp traffic

### 10.2.3 EDM Dataset

This dataset has been used for step four for the testing of the machine learning algorithms. This is a dataset of unlabeled netflow data gathered from the EDM network. It spans from the 18th of Februari, 2016 till the 24th of March 2016. The data spans from 10am till midnight. This dataset is used mainly to check whether the machine learning algorithms perform well on real-world data.

### 10.2.4 Cegeka Dataset

A dataset from Cegeka has been used. This dataset spans three days, from the third of April till the fifth of April. This dataset contains unlabeled netflow data. It also contains logs from firewalls. These logs can be used to check whether the machine learning algorithm has a good performance.

## 10.3 Feature selection

A lot of information can be found within a flow as explained in Section 9. The standard attributes of a flow have also been chosen to be used as features. The extra information such as the TCP flags that are set have not been used since they are not available in the datasets that have been used.

The country of origin has also been used in the implementation. However most services that provide such an API are limited in use and cannot be used to request data for millions of flows.

## 10.4 Algorithm selection

Both supervised and unsupervised algorithms have been used. The algorithms from Section 5 are the most common algorithms. Before more complex algorithms such as deep neural networks should be used, the more common and general algorithms should be tested.

### 10.4.1 Unsupervised learning

K-Means clustering is used in order to test whether results can be found using clustering algorithms. K-means is a simple clustering algorithm and already gives an indication whether a problem can be solved using clustering, or whether clustering offers no advantage.

One-class Support Vector machines are used in an attempt to use binary classification. They are quite fast in execution. They were used to find out whether it is a viable technique to preprocess incoming data and check whether a One-class Support Vector machine finds it to be abnormal behaviour before passing it to other algorithms.

### 10.4.2 Supervised learning

Support vector machines have been used in the implementation. It is a popular algorithm and can do both linear and non-linear classification which make it a promising choice to test in the implementation.

K-nearest Neighbors was the most promising algorithm. This algorithm is used extensively through throughout the implementation and the tests. The fact that the classification happens on basis of the different neighbors instead of trying to make a classifier seemed to fit the feature data better.

Neural networks also seemed very promising. However, the library *scikit-learn* does not implement neural networks and as such, they have not been used in the implementation.

Through the study of different machine learning algorithms, decision tree algorithms and Bayesian algorithms have also been discussed. They seemed less promising for the problem of intrusion detection. The difference between a normal flow and an abnormal flow is very slight and it seemed that these algorithms would make more mistakes. They are still used in the implementation.

# Chapter 11

## Evaluation

### 11.1 Supervised learning

#### 11.1.1 K-nearest Neighbors

#### 11.1.2 Support Vector machines

### 11.2 Unsupervised learning

#### 11.2.1 Clustering

#### 11.2.2 One-class Support Vector Machines

## Chapter 12

# Future work

### 12.1 Increasing accuracy

There are several possible methods that could be used to improve the accuracy of the implemented system. Currently the system generates alerts based on whether it thinks a flow is malicious or not. Incrementing the accuracy of the system could be solved by giving the system more data to learn from.

One method to be able to get more training data is to generate it. Using a Virtual Machine, different kinds of malware could be run. The generated network traffic can be used to train the system.

Another possible solution could be to implement packet analysis. A flow-based intrusion detection system processes a lot less data than a packet-based intrusion detection system. However, because IP flows contain less information than complete packets, a flow-based system cannot detect everything. It might be worth to find methods that incorporate both flow-based intrusion detection and packet-based intrusion detection.

One such method could be to use the IP flows to eliminate flows that are non-malicious and afterwards to perform packet analysis on the remaining flows. This may or may not have an impact on the accuracy and the different kinds of attacks that can be detected.

### 12.2 Intrusion Prevention Systems

It is useful to have an intrusion detection system. Intrusion prevention systems also have their place as a security tool. This thesis tried to find out whether machine learning techniques and IP flows can be used for an efficient intrusion detection system. Are these techniques also usable in the context of intrusion prevention?

There are two main issues with using a similar system for intrusion prevention. IP Flows are defined as an entire "flow", as a connection. At the moment a flow is received by the intrusion detection system, every packet has already been exchanged between the source and the destination. This means that "prevention" is more difficult since the intrusion has already happened. A solution to this could be to process incomplete flows. This means that the flow might still change and that communication is still ongoing. However, it is unknown how the system reacts to unknown flows.



The other issue is the performance of the machine learning algorithms. During the tests that were done in this thesis, no performance issues of the machine learning algorithms were noticed. But it is unknown how fast the machine learning algorithm needs to perform, should the intrusion detection system be fast enough.

### 12.3 Neural networks

As explained in Section 5.5, Neural networks are powerful machine learning algorithms. This is also suggested in other research such as [97]. They focussed on very specific methods on how neural networks can be used and did not use IP flows. However, neural networks have not been tested in this implementation.

## Chapter 13

# Conclusion

---

Write conclusion

### 13.1 IP Flows

### 13.2 Machine learning

# Bibliography

- [1] *Hospital declares 'internal state of emergency' after ransomware infection*, <http://krebsonsecurity.com/2016/03/hospital-declares-internet-state-of-emergency-after-ransomware-infection/>, Accessed: 2016-05-12.
- [2] *Sans: Host vs network based intrusion detection systems*, <https://cyber-defense.sans.org/resources/papers/gsec/host-vs-network-based-intrusion-detection-systems-102574>, Accessed: 2016-05-11.
- [3] M. J. Martin, *Smart grid: Cyber security*, <https://www.linkedin.com/pulse/smart-grid-cyber-security-martin-ma-mba-med-gdm-scpm-pmp>, Accessed: 2016-03-24.
- [4] R. G. Bace, *Intrusion Detection*. Indianapolis, IN, USA: Macmillan Publishing Co., Inc., 2000, ISBN: 1-57870-185-6.
- [5] A. Sundaram, "An introduction to intrusion detection", *Crossroads*, vol. 2, no. 4, pp. 3–7, Apr. 1996, ISSN: 1528-4972. DOI: 10.1145/332159.332161. [Online]. Available: <http://doi.acm.org/10.1145/332159.332161>.
- [6] H. Alaidaros, M. Mahmuddin, and A. Al Mazari, "An overview of flow-based and packet-based intrusion detection performance in high speed networks", in *Proceedings of the International Arab Conference on Information Technology*, 2011.
- [7] T.-L. Pao and P.-W. Wang, "Netflow based intrusion detection system", in *Networking, Sensing and Control, 2004 IEEE International Conference on*, IEEE, vol. 2, 2004, pp. 731–736.
- [8] *Sans: Understanding ips and ids: Using ips and ids together for defense in depth*, <https://www.sans.org/reading-room/whitepapers/detection/understanding-ips-ids-ips-ids-defense-in-depth-1381>, Accessed: 2016-05-11.
- [9] M. Golling, R. Hofstede, and R. Koch, "Towards multi-layered intrusion detection in high-speed networks", in *Cyber Conflict (CyCon 2014), 2014 6th International Conference On*, IEEE, 2014, pp. 191–206.
- [10] *Ips [intrusion prevention systems] – your 2nd line of defense*, <https://bastionnux.wordpress.com/2011/04/06/ips-intrusion-prevention-systems-your-2nd-line-of-defense/>, Accessed: 2016-05-11.
- [11] *Elprocus: Basic intrusion detection system*, <https://www.elprocus.com/basic-intrusion-detection-system/>, Accessed: 2016-05-12.
- [12] M. J. N. Jayveer Singh, "A survey on machine learning techniques for intrusion detection systems", *International Journal of Advanced Research in Computer and Communication Engineering*, vol. 2, no. 11, 2013.

- [13] A. Sperotto, G. Schaffrath, R. Sadre, C. Morariu, A. Pras, and B. Stiller, "An overview of ip flow-based intrusion detection.", *IEEE Communications Surveys and Tutorials*, vol. 12, no. 3, pp. 343–356, 2010. [Online]. Available: <http://dblp.uni-trier.de/db/journals/comsur/comsur12.html#SperottoSSMPS10>.
- [14] M. Uddin, R. Alsaqour, and M. Abdelhaq, "Intrusion detection system to detect ddos attack in gnutella hybrid p2p network", *Indian Journal of Science and Technology*, vol. 6, no. 2, pp. 4045–4057, 2013.
- [15] *Alienvault: Who we are*, <https://www.alienvault.com/who-we-are>, Accessed: 2016-05-11.
- [16] *Alienvault: Products*, <https://www.alienvault.com/products>, Accessed: 2016-05-11.
- [17] *Alienvault: Intrusion detection techniques: Methods and best practices*, <https://www.alienvault.com/blogs/security-essentials/intrusion-detection-techniques-methods-best-practices>, Accessed: 2016-05-11.
- [18] G. Kurundkar, N. Naik, and S. Khamitkar, "Network intrusion detection using snort", *International Journal of Engineering Research and Applications*, vol. 2, no. 2, pp. 1288–1296, 2012.
- [19] N. Duffield, P. Haffner, B. Krishnamurthy, and H. Ringberg, "Rule-based anomaly detection on ip flows", in *INFOCOM 2009, IEEE*, IEEE, 2009, pp. 424–432.
- [20] K. Veeramachaneni and I. Arnaldo, "Ai2: Training a big data machine to defend",
- [21] *Mit designed ai2, the system that can detect 85attacks*, <http://securityaffairs.co/wordpress/46480/security/ai2-system.html>, Accessed: 2016-05-12.
- [22] N Paulauskas and E Garsva, "Computer system attack classification", *Elektronika ir Elektrotechnika*, vol. 66, no. 2, pp. 84–87, 2015.
- [23] S. Hansman and R. Hunt, "A taxonomy of network and computer attacks", *Computers & Security*, vol. 24, no. 1, pp. 31–43, 2005.
- [24] *Man in the middle part 1: Introduction*, <https://toschprod.wordpress.com/2011/11/06/man-in-the-middle-intro/>, Accessed: 2016-04-28.
- [25] *A peek at the future of botnet evolution*, <http://securitywatch.pcmag.com/none/309491-a-peek-at-the-future-of-botnet-evolution>, Accessed: 2016-04-28.
- [26] A.-S. Kim, H.-J. Kong, S.-C. Hong, S.-H. Chung, and J. W. Hong, "A flow-based method for abnormal network traffic detection", in *Network operations and management symposium, 2004. NOMS 2004. IEEE/IFIP*, IEEE, vol. 1, 2004, pp. 599–612.
- [27] D. Moore, V. Paxson, S. Savage, C. Shannon, S. Staniford, and N. Weaver, "Inside the slammer worm", *IEEE Security & Privacy*, no. 4, pp. 33–39, 2003.

- [28] Y. Abuadlla, G. Kvascev, S. Gajin, and Z. Jovanovic, "Flow-based anomaly intrusion detection system using two neural network stages", *Computer Science and Information Systems*, vol. 11, no. 2, pp. 601–622, 2014.
- [29] Z. Zhu, G. Lu, Y. Chen, Z. J. Fu, P. Roberts, and K. Han, "Botnet research survey", in *Computer Software and Applications*, 2008. COMP-SAC'08. 32nd Annual IEEE International, IEEE, 2008, pp. 967–972.
- [30] Coursera machine learning stanford university, <https://www.coursera.org/learn/machine-learning>, Accessed: 2016-02-09.
- [31] Coursera machine learning lectures stanford university, <https://class.coursera.org/ml-003/lecture>, Accessed: 2016-02-09.
- [32] scikit-learn linear regression, [http://scikit-learn.org/stable/modules/linear\\_model.html](http://scikit-learn.org/stable/modules/linear_model.html), Accessed: 2016-02-15.
- [33] Stanford cs229 machine learning, <http://cs229.stanford.edu/materials.html>, Accessed: 2016-04-18.
- [34] Intuitive machine learning : Gradient descent simplified, <http://ucanalytics.com/blogs/intuitive-machine-learning-gradient-descent-simplified/>, Accessed: 2016-04-20.
- [35] Stackoverflow gradient descent convergence how to decide convergence?, <http://stackoverflow.com/questions/17289082/gradient-descent-convergence-how-to-decide-convergence>, Accessed: 2016-04-18.
- [36] Coursera machine learning gradient descent stanford university, <https://class.coursera.org/ml-003/lecture/11>, Accessed: 2016-02-09.
- [37] Cis 520 machine learning - 2015, <https://alliance.seas.upenn.edu/~cis520/dynamic/2014/wiki/index.php?n=Lectures.Classification>, Accessed: 2016-04-22.
- [38] Coursera machine learning cost function stanford university, <https://class.coursera.org/ml-003/lecture/6>, Accessed: 2016-02-09.
- [39] C. M. Bishop, "Pattern recognition and machine learning. springer".
- [40] Logistic regression, [http://www.holehouse.org/mlclass/06\\_Logistic\\_Regression.html](http://www.holehouse.org/mlclass/06_Logistic_Regression.html), Accessed: 2016-04-22.
- [41] Test-Driven Data Analysis in defence of xml: Exporting and analysing apple health data, <http://www.tdda.info/>, Accessed: 2016-04-18.
- [42] DTreg svm - support vector machines, <https://www.dtreg.com/solution/view/20>, Accessed: 2016-04-18.
- [43] Why is logistic regression a linear classifier?, <http://stats.stackexchange.com/questions/93569/why-is-logistic-regression-a-linear-classifier>, Accessed: 2016-04-24.
- [44] Linear versus nonlinear classifiers, <http://nlp.stanford.edu/IR-book/html/htmledition/linear-versus-nonlinear-classifiers-1.html>, Accessed: 2016-04-24.

- [45] *Kernel methods and nonlinear classification*, <https://www.cs.utah.edu/~piyush/teaching/15-9-slides.pdf>, Accessed: 2016-04-24.
- [46] *A tour of machine learning algorithms*, <http://machinelearningmastery.com/a-tour-of-machine-learning-algorithms/>, Accessed: 2016-03-11.
- [47] *Ricky Ho predictive analytics: Neuralnet, bayesian, svm, knn*, <http://horicky.blogspot.be/2012/06/predictive-analytics-neuralnet-bayesian.html>, Accessed: 2016-04-18.
- [48] *Wolfram mathworld: Distance*, <http://mathworld.wolfram.com/Distance.html>, Accessed: 2016-04-24.
- [49] *Creating my first algorithm from scratch – euclidean distance and pearson correlation*, <https://bigsnarf.wordpress.com/2012/03/25/creating-my-first-algorithm-from-scratch-euclidean-distance-and-pearson-correlation/>, Accessed: 2016-04-24.
- [50] *Hamming distance*, <https://people.rit.edu/rmb5229/320/project3/hamming.html>, Accessed: 2016-04-24.
- [51] *Chebyshev distance explanation*, <https://chessprogramming.wikispaces.com/Distance>, Accessed: 2016-04-25.
- [52] *Chebyshev distance*, [https://en.wikipedia.org/wiki/Chebyshev\\_distance](https://en.wikipedia.org/wiki/Chebyshev_distance), Accessed: 2016-04-25.
- [53] *Privacy preserving scheme for location-based services*, [http://file.scirp.org/Html/6-7800068\\_18792.htm](http://file.scirp.org/Html/6-7800068_18792.htm), Accessed: 2016-04-24.
- [54] R. V. Giuseppe Jurman Samantha Riccadonna and C. Furlanello, “Canberra distance on ranked lists”, Accessed: 2016-04-24. [Online]. Available: <http://mpba.fbk.eu/files/publications/jurman09canberra.pdf>.
- [55] *The minkowski distance (lp)*, [https://renatocamorim.com/2012/09/20/calculating-the-minkowski-center-1\\_p/](https://renatocamorim.com/2012/09/20/calculating-the-minkowski-center-1_p/), Accessed: 2016-04-25.
- [56] *Simplification and shift in cognition of political difference: Applying the geometric modeling to the analysis of semantic similarity judgment*. [https://openi.nlm.nih.gov/detailedresult.php?img=PMC3108968\\_pone.0020693.g001&req=4](https://openi.nlm.nih.gov/detailedresult.php?img=PMC3108968_pone.0020693.g001&req=4), Accessed: 2016-04-25.
- [57] *Scikit learn: 1.6. nearest neighbors*, <http://scikit-learn.org/stable/modules/neighbors.html>, Accessed: 2016-04-26.
- [58] “Intelligent data engineering and automated learning - ideal 2007: 8th international conference, birmingham, uk, december 16-19, 2007. proceedings”, in. Berlin, Heidelberg: Springer Berlin Heidelberg, 2007, ch. Using a Genetic Algorithm for Editing k-Nearest Neighbor Classifiers, pp. 1141–1150, ISBN: 978-3-540-77226-2. DOI: [10.1007/978-3-540-77226-2\\_114](https://doi.org/10.1007/978-3-540-77226-2_114). [Online]. Available: [http://dx.doi.org/10.1007/978-3-540-77226-2\\_114](http://dx.doi.org/10.1007/978-3-540-77226-2_114).
- [59] *K-means clustering*, <http://www.onmyphd.com/?p=k-means.clustering&ckattempt=1>, Accessed: 2016-04-26.

- [60] *A tutorial on clustering algorithms*, [http://home.deib.polimi.it/matteucc/Clustering/tutorial\\_html/kmeans.html](http://home.deib.polimi.it/matteucc/Clustering/tutorial_html/kmeans.html), Accessed: 2016-04-26.
- [61] *Python pattern recognition: K-means*, <http://pypr.sourceforge.net/kmeans.html>, Accessed: 2016-04-26.
- [62] *Scikit learn: One-class svm with non-linear kernel (rbf)*, [http://scikit-learn.org/stable/auto\\_examples/svm/plot\\_oneclass.html](http://scikit-learn.org/stable/auto_examples/svm/plot_oneclass.html), Accessed: 2016-04-26.
- [63] A. Ng, *Stanford machine learning: Neurons and the brain*, <https://class.coursera.org/ml-005/lecture/44>, Accessed: 2016-04-26.
- [64] —, *Stanford machine learning: Model representation II*, <https://class.coursera.org/ml-005/lecture/46>, Accessed: 2016-04-26.
- [65] *Stackoverflow what is a 'layer' in a neural network*, <http://stackoverflow.com/questions/35345191/what-is-a-layer-in-a-neural-network>, Accessed: 2016-04-18.
- [66] —, *Stanford machine learning: Multiclass classification*, <https://class.coursera.org/ml-005/lecture/49>, Accessed: 2016-04-26.
- [67] *Openstar cnx neural network implementation*, <https://cnx.org/contents/6CAkQax3@1/Neural-Networks>, Accessed: 2016-04-18.
- [68] —, *Stanford machine learning: Model representation II*, <https://class.coursera.org/ml-005/lecture/46>, Accessed: 2016-04-26.
- [69] —, *Stanford machine learning: Cost function*, <https://class.coursera.org/ml-005/lecture/50>, Accessed: 2016-04-26.
- [70] —, *Stanford machine learning: Backpropagation*, <https://class.coursera.org/ml-005/lecture/52>, Accessed: 2016-04-26.
- [71] —, *Stanford machine learning: Training a neural network*, <https://class.coursera.org/ml-005/lecture/56>, Accessed: 2016-04-25.
- [72] —, *Stanford machine learning: Autonomous driving*, <https://class.coursera.org/ml-005/lecture/57>, Accessed: 2016-04-26.
- [73] *Quadcopter navigation in the forest using deep neural networks*, <https://www.youtube.com/watch?v=umRdt3zGgpU>, Accessed: 2016-04-26.
- [74] A. C. Ian Goodfellow Yoshua Bengio, "Deep learning", Book in preparation for MIT Press, 2016, [Online]. Available: <http://www.deeplearningbook.org>.
- [75] *Scikit learn: 1.10. decision trees*, <http://scikit-learn.org/stable/modules/tree.html>, Accessed: 2016-04-26.
- [76] *Arjarapu blog space: Decision tree learning*, <http://www.arjarapu.com/wordpress/2012/04/decision-tree-learning/>, Accessed: 2016-04-26.



- [77] Wolfram mathworld: Bayes' theorem, <http://mathworld.wolfram.com/BayesTheorem.html>, Accessed: 2016-04-26.
- [78] Scikit learn: 1.9. naïve bayes, [http://scikit-learn.org/stable/modules/naive\\_bayes.html](http://scikit-learn.org/stable/modules/naive_bayes.html), Accessed: 2016-04-26.
- [79] Apriori algorithm, [https://en.wikipedia.org/wiki/Apriori\\_algorithm](https://en.wikipedia.org/wiki/Apriori_algorithm), Accessed: 2016-04-26.
- [80] R interface to oracle data mining, <http://www.analyticbridge.com/group/rprojectandotherfreesoftwaretools/forum/topics/r-interface-to-oracle-data>, Accessed: 2016-04-18.
- [81] A. Ng, *Stanford machine learning: Anomaly detection vs supervised learning*, <https://class.coursera.org/ml-005/lecture/93>, Accessed: 2016-04-26.
- [82] Normal distribution, [https://en.wikipedia.org/wiki/Normal\\_distribution](https://en.wikipedia.org/wiki/Normal_distribution), Accessed: 2016-04-26.
- [83] Online learning, [https://en.wikipedia.org/wiki/Online\\_machine\\_learning](https://en.wikipedia.org/wiki/Online_machine_learning), Accessed: 2016-04-26.
- [84] Improving the identification performance of an industrial process using multiple neural networks, <http://article.sapub.org/10.5923.j.ajis.20120204.02.html>, Accessed: 2016-04-26.
- [85] —, *Coursera machine learning, evaluating a hypothesis*, <https://www.coursera.org/learn/machine-learning/lecture/yfbJY/evaluating-a-hypothesis>, Accessed: 2016-04-22.
- [86] Model selection, [https://en.wikipedia.org/wiki/Model\\_selection](https://en.wikipedia.org/wiki/Model_selection), Accessed: 2016-04-22.
- [87] Cornell cs4670, [http://www.cs.cornell.edu/courses/cs4670/2015sp/lectures/lec34\\_datasets\\_web.pdf](http://www.cs.cornell.edu/courses/cs4670/2015sp/lectures/lec34_datasets_web.pdf), Accessed: 2016-04-18.
- [88] Which “mean” to use and when?, <http://stats.stackexchange.com/questions/23117/which-mean-to-use-and-when>, Accessed: 2016-04-22.
- [89] R. Sommer and V. Paxson, “Outside the closed world: On using machine learning for network intrusion detection”, in *Proceedings of the 2010 IEEE Symposium on Security and Privacy*, ser. SP '10, Washington, DC, USA: IEEE Computer Society, 2010, pp. 305–316, ISBN: 978-0-7695-4035-1. DOI: 10.1109/SP.2010.25. [Online]. Available: <http://dx.doi.org/10.1109/SP.2010.25>.
- [90] A. Subaira and P. Anitha, “Efficient classification mechanism for network intrusion detection system based on data mining techniques: A survey”, in *Intelligent Systems and Control (ISCO), 2014 IEEE 8th International Conference on*, IEEE, 2014, pp. 274–280.
- [91] O. Al-Jarrah, A. Siddiqui, M. Elsalamouny, P. D. Yoo, S. Muhaidat, and K. Kim, “Machine-learning-based feature selection techniques for large-scale network intrusion detection”, in *Distributed Computing Systems Workshops (ICDCSW), 2014 IEEE 34th International Conference on*, IEEE, 2014, pp. 177–181.
- [92] Netflow, <https://en.wikipedia.org/wiki/NetFlow>, Accessed: 2016-04-28.



- [93] *Ip api json format*, <http://ip-api.com/docs/api:json>, Accessed: 2016-04-28.
- [94] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, "Scikit-learn: Machine learning in Python", *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [95] S. García, M. Grill, J. Stiborek, and A. Zunino, "An empirical comparison of botnet detection methods", *Computers & security*, vol. 45, pp. 100–123, 2014. [Online]. Available: <http://mcfp.weebly.com/the-ctu-13-dataset-a-labeled-dataset-with-botnet-normal-and-background-traffic.html>.
- [96] A. Sperotto, R. Sadre, D. F. van Vliet, and A. Pras, "A labeled data set for flow-based intrusion detection", in *Proceedings of the 9th IEEE International Workshop on IP Operations and Management, IPOM 2009, Venice, Italy*, ser. Lecture Notes in Computer Science, vol. 5843, Springer Verlag, 2009, pp. 39–50. [Online]. Available: <https://traces.simpleweb.org/traces/netflow/netflow2/>.
- [97] R. Koch and G. Dreo, "Fast learning neural network intrusion detection system", in *Scalability of Networks and Services*, Springer, 2009, pp. 187–190.

# Appendix A

## User guide

### A.1 Running

The main program can be run using the command in Listing A.1. If no config file is passed as argument the program will look for a config file named "config.json" in the current directory.

```
python main.py [config-file-location]
```

### A.2 Running multiple tests

It is also possible to run multiple tests automatically. This can be done using the command as seen in Listing A.2. If no config file is passed as argument the program will look for a config file named "testing.json" in the current directory. The layout of the config file is fairly simple. It should be a list of config files that have to be run as seen in Listing A.2.

```
python testing.py [config-file-location]
```

```
1 [
2     "main/config.json"
3 ]
```

### A.3 Config file

The config file contains the main variables that define the execution of the program. If the format of an attribute is not correct, the program either skips that value or halts program execution.

There are five important boolean settings in the config file. **enable-IDS** is a boolean which specifies whether the IDS should run or not. **pcap-to-flow** is whether the small, built-in flow-converter should be used. This is used to convert *pcap* files to flow files. **print-labels** activates a section of the program which extracts and prints the labels that belong to a given dataset. This could be useful when trying to find out which labels the currently selected dataset will use.

The booleans **check** and **sniffer** can only be used when **enable-IDS** is set

to *True*. **check** is used to activate or disable the checking of datasets after the training phase. **sniffer** is used to activate the built-in packet sniffer. The packet sniffer sniffs packets and builds flows, which are then fed to the machine learning algorithm.

```

1 {
2     "enable-IDS" : true,
3     "pcap-to-flow" : false,
4     "print-labels" : false
5
6     "check" : true,
7     "sniffer" : false,
```

**pcap-files** is a list of (source, destination) tuples for files that have to be converted using the built-in flow converter.

```

1     "pcap-files" : [
2         {
3             "src" : "../..//test/dosattack.pcap",
4             "dest" : "../..//test/dosattack2.flow"
5         }
6     ],
```

**algorithm** is used to define which algorithm is used. The value of this attribute should match the name of one of the implemented machine learning algorithm classes. **featureClass** is used to select which class to use to retrieve features from a flow. **trainer** is used to select the training method. When the value is set to "Trainer", it is possible to select the training class in the data sets.

```

1     "algorithm" : "KNeighborsClassifier",
2     "featureClass" : "FlowFeature",
3     "trainer" : "Trainer",
```

Some attributes in the config file are used to save and use the models that the machine learning algorithms make. **model\_dir** is used to define the directory location. **model** defines the name of the model file. **use\_model** and **store\_model** are respectively used to activate usage and saving of the model.

```

1     "model_dir" : "../..//models/",
2     "model" : "model.mdl",
3     "use_model" : false,
4     "store_model" : true,
```

**data-sets** indicates the data to be used when training. The layout of each

item in the list is dependent of which loader should be used. A netflow file has the attributes: **file**, **from** and **to**. An SQL trainer is also implemented. Here the attributes should indicate the **host**, the **user**, the **database** and the **password**. The **type** can be used to indicate which type of data it is. This should be the name of a trainer class.

The **check-sets** are used to check the performance of the machine learning algorithm. Here the **type** indicates which type the data is. The possible values are: "file" and "sql".

```

1  "data-sets" : [
2      {
3          "file" : "../test/capture20110815.
              binetflow",
4          "from" : 0,
5          "to" : 40000
6      },
7      {
8          "host" : "localhost",
9          "type" : "SQLTrainer",
10         "user" : "user",
11         "db" : "dataset",
12         "password" : "password"
13     }
14 ],
15
16
17 "check-sets" : [
18     {
19         "file" : "../test/capture20110815.
                binetflow",
20         "type" : "file",
21         "from" : 0,
22         "to" : -1
23     }
24 ],

```

**feature-file** is a file that can contain information to be used in the feature class. The layout of this file is dependent on which "featureClass" is used. **logger** indicates the location of the log file to use. **good-labels** indicate which labels are considered to belong to normal behaviour.

```

1  "feature-file" : "features.json",
2  "logger" : "log.txt",
3  "good-labels" : "good.txt",

```

The following attributes are used in the small sniffer that was implemented. **protocol-file** is a file that selects which protocols to use. **timeout** is a value that define the amount of milliseconds to wait before closing a flow after the last packet in that flow.

```
1   "protocol-file" : "protocols.json",  
2   "timeout" : 5000,  
3 }
```

## Appendix B

# Developer Guide

The implementation is made to be modular. Adding new algorithms or features is as simple as adding a new class.

### B.1 Machine learning module

New machine learning algorithms can be added to the implementation by creating a new class and placing the class in the *ml.py* file. A base class, **MLAlgorithm**, has to be inherited from and two methods have to be implemented:

---

```
class SomeAlgorithm(MLAlgorithm):
    def train(self, data_set, target_set=None):
        data_set is an array of features
        target_set is an array of labels
        (or none for unsupervised learning)

    def predict(self, sample, corr=None):
        sample is a feature that is going
            to be used for prediction
        corr is the correct label that
            should be predicted
```

---

### B.2 Feature module

A feature class is used to parse a flow. The **Flow** class is a simple "struct" like class that contains the necessary member variables for a flow. The **Flows** class is a collection of flows. To make a new featureClass, a base class, **BasicFeature**, has to be inherited from and two methods have to be implemented:

---

```
class FlowFeature(BasicFeature):
    def __init__(self, file):
        file is a json file that may contain extra
            information to make a featureClass
            more reusable

    def make_record(self, flow):
        flow is an instance of the Flow class
```

---

---

```
return an array of numbers to be fed
        to a machine learning algorithm
```

---

### B.3 Loader module

The loader module is responsible for loading the flows from any external source. The external source could be a text file or a SQL database. The new class should be placed in the *loader.py* file. The loader class should be defined using a *load* and a *get\_netflow* method. The parameters of *load* can be chosen at will since the loading has to be implemented in a training class.

---

```
class LoaderName:
    def load(self, ...):
        load the data from the source

    def get_netflow(self):
        return the netflow
```

---

### B.4 Prediction Module

If a new type of loader needs to be used in the check-sets, a class has to be added to the Prediction module in the *prediction\_type.py* file. A base class, **PredictionLoader**, has to be inherited from and a method has to be implemented: "load(*self*, *data*)". "data" is a dictionary loaded from the config file.

---

```
class PredictionFile(PredictionLoader):
    def load(self, data):
        return the netflow
```

---

### B.5 Training module

The training module is responsible for all training classes. New types of training can be added to the implementation by creating a new class and placing the class in the *train.py* file. A base class, **Trainer**, has to be inherited from and a method has to be implemented. The method has several parameters. "algorithm" is an instance of a machine learning algorithm. "data" is a dictionary loaded from the config file from data-sets. "feature" is an instance of the selected featureClass. "good\_labels" is a list of all labels that are considered normal behaviour.

*self.default* is a method implemented in the Trainer class. This method loads the file using the specified loader and file. Afterwards it extracts the features and feeds them to the machine learning algorithm.

---

```
class TrainerName(Trainer):
```

---

---

```
def train(self, algorithm, data, feature,
          good_labels):
    load the data
    return self.default(algorithm, loader,
                        feature, file)
```

---

## B.6 Results module

The results module is the biggest module and is constructed from three files. The first file, *logger.py*, contains the main **Logger** class. These methods write to the log file that is specified in the config file.

*visualise.py* contains all methods that are used to visualise the machine learning algorithms. These visualisation methods can be injected into the program execution by using them in a machine learning class. This is required since the visualisation methods depend highly on the "train" and "predict" methods of the machine learning algorithms.

Finally, the *results.py* file receives all output from the prediction algorithms. They can filter out predictions that are considered normal behaviour and evaluate the machine learning algorithm using the amount of false positives and negatives.

## B.7 Other

*main.py* and *testing.py* are the main entry-points of the program. *main.py* puts together all different components to make the program execute. The *config.py* file is responsible for parsing the config file. Any new additions to the config file need to be reflected in this file. *predictor.py* is the file responsible for running the different check-sets. *sniffer.py* contains the small implementation of the sniffer, the real-time intrusion detection implementation.



# Appendix C

## Meetings

### C.1 Meeting 1: 09 Feb 2016

aanwezig: Peter Quax, Bram Bonne, Pieter Robyns, Axel Faes

Dit is de eerste bijeenkomst met de begeleiders en promotor. Er is dus geen rapportering mogelijk van een vorige bijeenkomst. Tijdens de bijeenkomst is beslist om een *intruder detection system* te bestuderen en te implementeren.

De actiepunten die gedaan moeten worden:

- Beslissen voor wie het systeem gemaakt moet worden. Gaat dit voor end users zijn, of voor grote data centers. Hieraan hangt vast welke data (packets of netflow) gebruikt moet worden.
- Bekijken hoe machine learning algoritmes gebruikt kunnen worden in een *intruder detection system*.
- Bekijken wat netflow is.
- Er moet gekeken worden naar de manier waarop anomalies gegenereerd gaan worden om het systeem te testen/trainen.

Volgende afspraken zijn gemaakt:

- Er is gevraagd om te zorgen dat het systeem ook op correcte wijze informatie kan weergeven aan gebruikers. Tijdens het semester moet bekeken worden hoe deze weergave moet gebeuren.
- Libraries gebruiken indien mogelijk, om te vermijden dat het wiel opnieuw uitgevonden word.
- Er is de mogelijkheid geboden om aan de thesis te werken op het EDM.
- Er is afgesproken om *Overleaf* te gebruiken om de thesis in te schrijven.
- Een ruwe planning voor het werk moet gemaakt worden tegen 12 Feb.
- Een wekelijkse meeting is vastgelegd. Dit om 10:00 elke vrijdag.
- Begin mei moet een eerst draft van de thesis klaar zijn en eind mei moet de finale draft af zijn.
- Er moet een vulgariserende tekst gemaakt worden en een postersessie gegeven worden (op 29 juni).

## C.2 Meeting 2: 12 Feb 2016

aanwezigen: Bram Bonne, Pieter Robyns, Axel Faes

Dit is de tweede bijeenkomst met mijn begeleider. Netflow bevat op zichzelf niet zoveel informatie, maar het is toch handig om te kijken welke bevindingen gemaakt kunnen worden met deze data. Mogelijks kan er, indien gevonden wordt dat netflow alleen niet genoeg informatie bevat, ook gebruikt gemaakt worden van packet data.

Er is de mogelijkheid besproken om eventueel meerdere machine learning algoritmes te implementeren en te bekijken in welke situaties welke algoritmes beter werken.

De actiepunten die gedaan zijn:

- *Beslissen voor wie het systeem gemaakt moet worden.*: Dit gaat gedaan worden voor data centers
- Er zijn verschillende classificaties van machine learning algoritmes gevonden die gebruikt kunnen worden.
- Verschillende grote data sets van netflow en packets met sporen van anomalies zijn gevonden. Alsook programma's om verkeer te genereren.

Volgende actiepunten zijn besproken:

- Verder uitwerken van welke machine learning algoritmes gebruikt kunnen worden
- Bekijken netflow v9

## C.3 Meeting 3: 19 Feb 2016

aanwezigen: Bram Bonne, Pieter Robyns, Axel Faes

Professor Quax is aan het bekijken ofdat ik (gelabelde) netflow data kan verkrijgen van Cegeka. Dit zou heel handig zijn om mijn implementatie te testen op real world data.

Voorlopig moet ik enkel focussen op een passive intrusion detection systeem, geen preventie en niet direct inline in het netwerkverkeer. Ook de visualisatie moet later bekeken worden, de gebruiker is een netwerkadministrator. Er is tevens besproken dat python zelf mogelijks te traag is om packet sniffing op een goede snelheid uit te voeren. Hiervoor zou ik wireshark kunnen gebruiken (of de command line versie). Er is besproken om eventueel zelf datasets te genereren door malware te runnen op een VM of aparte machine.

De datastructuur voor de machine learning algoritmes is bekeken. Ik moet eens bekijken hoe de timestamps van de flowdata gebruikt kunnen worden. Om de effectiviteit (van de machine learning algoritmes) mogelijks te

verhogen ga ik eens bekijken of ip-adressen ingedeeld kunnen worden in country-of-origin of iets dergelijks. Dit zou de machine learning algoritmes de mogelijkheid bieden om ook op deze parameter te bekijken of data malicious is of niet.

De actiepunten die gedaan zijn:

- Er is al een basis implementatie uitgewerkt voor het IDS
- De netflow structuur is bekeken en er is een datastructuur opgesteld die gefeed kan worden aan verschillende machine learning algoritmes.
- Progressie in de machine learning cursus: chapter 3 van de 18.

Volgende actiepunten zijn besproken:

- Beginnen aan de thesis: het schrijven van een hoofdstuk over machine learning en over hoe deze algoritmes toegepast kunnen worden op een intrusion detection systeem.
- Verder werken in de machine learning cursus.
- Ik moet eens bekijken ofdat ik een programma vind om pcap files om te zetten naar netflow. Anders moet ik dit zelf schrijven.

Ik heb ook een korte planning gemaakt van hoe de thesis eruit zou zien:

- Inleiding:
  - wat is een IDS
  - Waarom is er gekozen voor dit type IDS (host vs netwerk)
  - Waarom voor data centers
  - Waarom netflow
  - Waarom machine learning
- Wat is machine learning
- Hoe passen we machine learning toe op IDE en wat zijn de voor/-naden
- Welke machine learning algortimes zijn wel/niet gebruikt
- Wat zijn de voor/nadelen van netflow
- Hoe met combinatie netflow/packets (Als dit gedaan zou worden)
- Welke data sets zijn gebruikt
- Wat zijn de bevindingen
- Hoe kan visualisatie/feedback gebeuren (richting admin en richting automatische preventie)
- Conclusie

## C.4 Meeting 4: 26 Feb 2016

aanwezig: Bram Bonne, Axel Faes

Deze week is voornamelijk besteed aan de implementatie. Er is een netflow exporter geschreven. Er is bekeken ofdat timestamps gebruikt kunnen worden en ofdat ip-adressen opgedeeld kunnen worden per land. Er is besloten dat dit zeer weinig effect heeft op de accuraatheid van de machine learning algoritmes.

Momenteel zijn Support vector machines en K-nearest Neighbor Classifier algoritmes bekeken. Het K-nearest Neighbor Classifier algoritme is zeer efficiënt (98%).

In een later stadium kan bekeken worden om eventueel verdere analyse te doen op de data die malicious gevonden is, eventueel door pakketten te analyseren, of nogmaals door machine learning technieken. Er kan ook eens bekeken worden om een VM op te zetten, en daarin malware te runnen en dit verkeer te monitoren. Hierbij zouden eigen datasets gegenereerd kunnen worden.

De machine learning cursus is gevolgd tot hoofdstuk 7. De cursus zou normaal af moeten zijn binnen 2 weken.

De actiepunten die gedaan zijn:

- Er is al een netflow exporter geschreven
- Er zijn experimenten uitgevoerd m.b.t de datastructuur die meegegeven wordt aan de machine learning cursus.
- Progressie in de machine learning cursus: chapter 7 van de 18.
- Er is begonnen aan de thesis.
- Het zou interessant zijn om eens te kijken ofdat ip-adressen opgedeeld kunnen worden in subnets.

Volgende actiepunten zijn besproken:

- Focussen op de thesis
- Verder werken in de machine learning cursus.

## C.5 Meeting 5: 04 Mar 2016

aanwezig: Bram Bonne, Axel Faes

Deze week is voornamelijk besteed aan de thesis en aan het leren van de machine learning cursus. De machine learning cursus is gevolgd tot hoofdstuk 10. De cursus zou tegen volgende meeting af moeten zijn. De algemene

structuur van de thesistekst is nagekeken. Het hoofdstuk over "Attack classification" moet uitgebreid worden met een algemene uitleg over hoe aanvallen gedetecteerd kunnen worden. Het hoofdstuk over de gebruikte data-sets moet samengevoegd worden met het hoofdstuk dat de implementatie beschrijft. Het hoofdstuk over voor/nadelen van machine learning voor intrusion detection systemen moet verwerkt worden in het algemene hoofdstuk over machine learning.

De actiepunten die gedaan zijn:

- Verder werken aan ML cursus
- Schrijven aan thesis.

Volgende actiepunten zijn besproken:

- Verder werken aan ML cursus
- Schrijven aan thesis.

## C.6 Tussentijdse presentatie: 08 Mar 2016

aanwezigen: Maarten Wijnants, Peter Quax, Wim Lamotte, Jori Liesenborgs, Wouter vanmontfort, Pieter Robyns, Robin Marx, Bram Bonne, Axel Faes

Er is een tussentijdse presentatie geweest waarbij ik mijn huidige progressie moest tonen en een planning moest geven. De presentatie zelf is goed verlopen. Na de presentatie heb ik verschillende vragen gekregen.

Veel vragen die gesteld waren, waren bedoeld om te kijken of we het nut/-doel van de bachelorthesis kennen en hoe we de invulling correct doen. Ook een belangrijk aspect is hoe het valideren van de correctheid van de experimenten die gedaan zijn/worden zal gebeuren.

Een opmerking was dat ik ook bestaande Intrusion detection systemen moet bekijken en ofdat deze machine learning gebruiken. Dan is ook belangrijk waarom ze het wel of niet gebruiken.

Er was verwacht dat ik al iets verder stond met de Machine learning cursus. Hierdoor kon ik niet altijd op de volledige diepgang de gestelde vragen beantwoorden. Ik begreep ook niet altijd de onderliggende vraag waardoor ik te oppervlakkig antwoorde. Qua machine learning algoritmes moest ik goed opletten voor overfitting en uitleggen hoe ik hiermee omga.

Er zijn ook vragen gesteld m.b.t mijn geplande extra om het intrusion detection systeem real-time te maken. Normaal wordt een flow pas doorgegeven als deze volledig afgesloten is, een mogelijke piste zou zijn om flows al te bekijken ook al zijn ze nog niet afgesloten. Ook het runnen van een VM met malware erop om zelf data-sets te generen is bevestigd dat een goed idee zou zijn. Als ik hiervoor infrastructuur nodig heb moet ik dit vragen.

Ik moet opletten met aanvallen die maar zeer weinig netwerktraffiek genereren. Ook moet ik goed beschrijven welke aanvallen wel of niet gedetecteert kunnen worden en uitleggen waarom. Dit staat momenteel al beschreven in mijn thesistekst. Ik zou ook mogelijkheden kunnen uitleggen die ervoor zouden kunnen zorgen dat ik toch alle (of een groot deel) van de aanvallen zou kunnen detecteren. Dit zou bv kunnen door toch packet-data te gaan bekijken.

Een algemene opmerking die gegeven was, was dat de presentatie visueler mocht zijn. Figuren en afbeeldingen zijn aangenamer om te tonen aan een publiek. Bij de postersessie moet er ook goed opgelet worden dat ik van persoon tot persoon bekijk hoe diep ik de materie uit mijn bachelorthesis kan uitleggen.

## C.7 Meeting 6: 11 Mar 2016

aanwezigen: Bram Bonne, Axel Faes

Deze week is voornamelijk besteed aan de thesis en het afmaken van de machine learning cursus. De machine learning cursus is volgens planning afgemaakt. De thesistekst moet ingestuurd worden zodat deze al nagekeken kan worden. Komende weken zal gespendeerd worden aan de implementatie en het uitzoeken hoe de flowdata gebruikt kan worden in de machine learning algoritmes.

De actiepunten die gedaan zijn:

- Afmaken Machine learning cursus
- Schrijven aan thesis.

Volgende actiepunten zijn besproken:

- Verdere implementatie algoritmes
- Bekijken hoe de flowdata gefeerd kan worden aan de machine learning algoritmes

## C.8 Meeting 7: 18 Mar 2016

aanwezigen: Bram Bonne, Pieter Robyns, Axel Faes

Deze week is voornamelijk besteed aan implementatie. Er zijn verschillende algoritmes geïmplementeerd. Deze algoritmes zijn K-Nearest Neighbors, K-Means, Lineair Kernel Support Vector Machines met One vs All classification, Lineair Kernel Support Vector Machines met One vs One classification en Gaussian Kernel Support Vector Machines.

Er zijn verschillende experimenten uitgevoerd. De poort nummers zijn geïmplementeerd in 2 varianten. De eerste variant splitst de poorten in categorieën (veel gebruikte poort nummers en niet-veel gebruikte poort nummers) als een binaire feature. De andere variant maakt van elk poort nummer een nieuwe binaire feature (die stelt of de poort gebruikt is of niet).

Deze variant geeft echter heel veel features, dit is enorm inefficient.

De gestelde feedback was om te kijken hoe goed het werkt als poort nummers als een continue feature voorgesteld wordt en om de categorieën op te splitsen in  $>1024$  en  $<1024$ . De IP-data kan opgesplitst worden in aparte features voor IPv6, IPv4 en MAC. Deze data kan mogelijks voorgesteld worden als continue data. Er is voorgesteld om ook andere algoritmes te implementeren. Om gebruik te kunnen maken van datasets van Cegeka moet ik een presentatie maken en een afspraak maken met professor Quax.

Er is feedback gegeven op de thesistekst. Ik moet als eerste goed letten op spelfouten. De inleiding moet algemener uitgelegd worden. Ook niet gebruikte concepter zoals Signature-based IDS moet dieper uitgelegd worden. De attack classification maakt al teveel assumpties over wat er gedetecteerd kan worden. Dit zou eerst algemener uitgelegd moeten worden. Het machine learning hoofdstuk bevat in het algemeen te weinig high-level beschrijvingen.

De actiepunten die gedaan zijn:

- Begin van hoe flowdata gebruikt kan worden
- Implementatie van K-Nearest Neighbors
- Implementatie van K-Means
- Implementatie van Lineair Kernel Support Vector Machines met One vs All classification
- Implementatie van Lineair Kernel Support Vector Machines met One vs One classification
- Implementatie van Gaussian Kernel Support Vector Machines.

Volgende actiepunten zijn besproken:

- Herschrijven en verwerken van feedback op de thesistekst
- Verdere implementatie: andere algoritmes
- Verdere implementatie: Ports indelen in  $>1024$  en  $<1024$
- Verdere implementatie: Ports indelen als continue data
- Verdere implementatie: IP indelen als continue data
- Verdere implementatie: Starttime instellen als unix time
- Maken presentatie voor Cegeka data set

## C.9 Meeting 8: 24 Mar 2016

aanwezig: Bram Bonne, Pieter Robyns, Axel Faes

Deze week is voornamelijk besteed aan implementatie. Er zijn verschillende experimenten uitgevoerd. De poort nummers zijn geïmplementeerd

zodanig dat er een indeling is tussen normale poorten ( $<1024$ ) en speciale poorten ( $>1024$ ). Er is ook gekeken ofdat poort data niet continu voorgesteld kan worden. Dit geeft echter slechtere resultaten dan te werken met de binaire indeling.

De IP-data kan opgesplitst worden in aparte features voor IPv6, IPv4 en MAC. Deze data is voorgesteld als continue data. Het is moeilijk om IP-adressen zelf te gaan onderverdelen in categorieën of subnets. Het indelen als continue data geeft dan ook de beste accuraatheid. De starttijd is ook geïmplementeerd om te voegen als feature. Dit gebeurt door te bekijken in welke dag van de week/ uur van de dag en minuut van het uur de data gegenereerd word. Echter had deze feature geen goede invloed op de accuraatheid

Er zijn enkele algoritmes verder geïmplementeerd zoals een Decision Tree learner en een Naive Bayes algoritmes. Beide zijn supervised learning technieken. De Naive Bayes had nog goede accuraatheid, de Decision Tree Learner gaf slechtere resultaten.

De actiepunten die gedaan zijn:

- Verdere implementatie: andere algoritmes
- Verdere implementatie: Ports indelen in  $>1024$  en  $<1024$
- Verdere implementatie: Ports indelen als continue data
- Verdere implementatie: IP indelen als continue data
- Verdere implementatie: Starttime instellen als unix time

Volgende actiepunten zijn besproken:

- Herschrijven en verwerken van feedback op de thesistekst
- Testen van de implementatie
- Bekijken van Unsupervised learning algoritmes

## C.10 Meeting 9: 01 April 2016

aanwezig: Bram Bonne, Peter Quax, Axel Faes

Deze week is voornamelijk besteed aan implementatie. Er is een nieuwe dataset gevonden. Deze dataset is afkomstig van de universiteit van Twente. Er was een honeypot opgesteld op het netwerk, alle data die hiermee gevangen is, is geëncrypteerd en een dataset mee gemaakt. De dataset bevat voornamelijk externe aanvallen.

Er is ook gefocused op het trachten te gebruiken van unsupervised learning algoritmes. Echter heeft dit weinig opgebracht. Er kan wel op een accurate manier onderscheidt gemaakt worden tussen malicious en niet malicious data, maar het is moeilijk om vervolgens af te leiden over welk type malicious data het gaat.



De features die gebruikt worden in de machine learning algoritmes zijn nog eens overlopen. Professor Quax kwam met het idee om IP-adressen eventueel op te delen in origine (zoals land). Dit kan gebeuren via services zoals WhoIs.

De EDM dataet kan afgehaald worden bij het kantoor van professor Quax. Er moet ook zo snel mogelijk een meeting georganiseerd worden met Cegeka.

De actiepunten die gedaan zijn:

- Herschrijven en verwerken van feedback op de thesistekst
- Testen van de implementatie
- Bekijken van Unsupervised learning algoritmes

Volgende actiepunten zijn besproken:

- Verwerken data EDM
- Implementatie van WhoIs als feature
- Maken presentatie voor Cegeka data set

## C.11 Meeting Cegeka: 06 April 2016

aanwezig: Peter Quax, Axel Faes, Cegeka

Er is een meeting geweest met Cegeka om de mogelijkheid te bespreken om data te kunnen gebruiken die afkomstig is van Cegeka, alsook om informatie te krijgen van de huidige IDS' die gebruikt worden.

In het begin is een korte presentatie gegeven waarin de eigenschappen van het te ontwikkelen systeem uitgelegd worden. Er wordt ook beschreven wat de algemene onderzoeksvragen zijn en hoe data nu gebruikt kan worden in machine learning algoritmes.

Een eerste vraag die gesteld is, is welke classificatie van 'onverwachte trafiek' er momenteel gebeurt in datacenter/hosting context. Cegeka werkt heel gelaagd. Voor de routers staat een DDoS protection systeem. Na de router staan zowel firewalls als SIEM devices. Voor grotere klanten is er extra beveiliging voorzien in de vorm van afgestelde IPS systemen en meer gedetailleerde threat detection. Zoveel mogelijk data wordt verwerkt, zowel flows als meer gedetailleerd. De systemen werken voornamelijk met een signature database en verwerken de data voornamelijk automatisch. De systemen moeten wel manueel afgesteld en onderhouden worden.

De classificatie gebeurt heel gedetailleerd door deze verschillende lagen. Er werd wel gesteld dat het veel interessanter is om outbound verkeer na te kijken in vergelijking met binnengaand verkeer. Zaken zoals port-scans zijn interessant om te weten maar gebeuren heel veel en kunnen al goed gedetecteerd worden.

Er was veel interesse naar een intrusion detectie systeem dat op basis van

netflow en machine learning technieken werkt. Er accurate van rond de 70-80 procent zou gezien worden als een goede accurate. Ze hebben ook liever false positives dan false negatives. Te veel alerts genereren is heel vervelend, maar het is belangrijker dat voldoende anomalieën gedetecteerd worden.

Er is gevraagd of het mogelijk is om data te verkrijgen. Dit was zeker mogelijk. De netflow en corresponderende logs van 3 dagen wordt geleverd. De logs worden zowel in binair als text formaat geleverd. Het binair formaat kan ingelezen door een programma dat de logs visualiseert. Dit programma wordt ook meegeleverd. Mogelijks zou ook output van het DDoS systeem geleverd kunnen worden. Om te bekijken of dat een flow gezien is als malicious of niet moet dit bekeken worden of deze flow voorkomt in de logs.

## C.12 Meeting 10: 12 April 2016

aanwezig: Bram Bonne, Axel Faes

Afgelopen week, woensdag 6 april, is de meeting geweest met Cegeka. Van deze meeting is een verslag gemaakt. Op maandag 4 april is de dataset van het EDM verkregen. Deze dataset bevat netflow data van het EDM netwerk van 18 februari tot 24 maart. Elke dag is netflow beschikbaar die voorgekomen is tussen 10u tot 24u. Deze zijn per 15 minuten gelogd in een file.

Er is kort gewerkt aan het verwerken van de data van het EDM. Dit gebeurt door de data te laten verwerken door verschillende algoritmes. De data waarvan vervolgens gesteld kan worden dat deze daadwerkelijk malicious is, zal doorgegeven worden aan professor Quax. Op het moment zijn er nog niet genoeg testen uitgevoerd om iets te kunnen zeggen over de data.

Er is geïmplementeerd dat de country-of-origin van een IP ook gebruikt kan worden als feature voor de machine learning algoritmes. Er is gevonden dat dit voornamelijk werkt voor data die van buitenaf komt, zoals port scans. Er is ook kort al geprobeerd om visualisaties te maken van de machine learning algoritmes. Zulke visualisaties zouden interessant kunnen zijn om ook in de thesistekst te gebruiken. Tegen volgende bijeenkomst moet voornamelijk gewerkt worden aan de thesistekst. De actiepunten die gedaan zijn:

- Meeting Cegeka
- Implementatie van WhoIs als feature
- Implementatie van kleine visualisatie van algoritmes
- Verkrijgen data van EDM

Volgende actiepunten zijn besproken:

- Herschrijven en verwerken van feedback op de thesistekst

- Visualisaties van machine learning algoritmes zijn interessant voor thesistekst.

### C.13 Meeting 11: 18 April 2016

aanwezig: Bram Bonne, Axel Faes

Ik heb mijn huidige vooruitgang laten zien van de bachelorproef tekst. Hierbij is een korte herschikking gekomen van de hoofdstukken. Ik moet het "Implementatie" hoofdstuk opsplitsen naar een hoofdstuk "Implementatie" en een hoofdstuk "Evaluatie". In "Implementatie" moet mijn implementatie zelf beschreven staan, in "Evaluatie" moeten de resultaten beschreven worden.

Verder is kort uitgelegd wat de algemene structuur is van het machine learning hoofdstuk en welke aanpassingen er gebeurd zijn. Het hoofdstuk is opgesplitst in meerdere delen zodanig dat de uitleg, de algoritmes en validatie in aparte hoofdstukken uitgelegd wordt. Het hoofdstuk "Preventie" moet ik aanpassen naar iets gelijkaardigs aan "Future work". Het hoofdstuk "Visualisatie" moet samengevoegd worden met "Implementatie". De actiepunten die gedaan zijn:

- Verwerken feedback op machine learning hoofdstukken.

Volgende actiepunten zijn besproken:

- Herschrijven en verwerken van feedback op de thesistekst
- Halen data bij Cegeka

### C.14 Meeting 12: 22 April 2016

aanwezig: Bram Bonne, Axel Faes

Ik heb mijn huidige vooruitgang laten zien van de bachelorproef tekst. De meeste tekst is al stukken beter. Er mogen bij enkele stukken nog numerieke voorbeelden komen te staan. Deze stukken zijn de regularisatie en de feature scaling. Het hoofdstuk over "Algorithms" moet nog herschreven worden. Tevens ga ik een "Summary" schrijven op het einde van het machine learning hoofdstuk zodanig dat nog snel en kort een samenvatting gegeven wordt.

De actiepunten die gedaan zijn:

- Verwerken feedback op machine learning hoofdstukken.

Volgende actiepunten zijn besproken:

- Verwerken feedback en herschrijven "Algorithms" tegen zondag 24 April
- Schrijven van tekst

## C.15 Meeting 13: 29 April 2016

aanwezig: Bram Bonne, Axel Faes

Het herschrijven van het "Algorithms" hoofdstuk heeft langer geduurd dan gedacht en is pas ingeleverd op woensdag 27 april 2016. Tijdens de meeting is toegelicht welke aanpassingen gebeurt zijn aan de thesis. Er is kort besproken dat er ook focus gelegd moet worden over het verwerken van de data van Cegeka. Dit zou al belangrijk zijn voor de eerste draft van de thesis.

De actiepunten die gedaan zijn:

- Beginnen met verwerken van data van Cegeka
- Schrijven korte summary op het einde van het machine learning hoofdstuk
- Numerieke voorbeelden plaatsen bij regularisatie en feature scaling
- Hoofdstuk "Algorithms" herwerken

Volgende actiepunten zijn besproken:

- Schrijven hoofdstuk IP Flows
- Herwerken van hoofdstuk "Machine learning for an IDS"
- Verder verwerken dataset Cegeka

## C.16 Meeting 14: 04 Mei 2016

aanwezig: Bram Bonne, Axel Faes

Er is feedback gegeven over de machine learning hoofdstukken. De inleiding is goed geschreven. Bij het gedeelte van Feature scaling moet ik 2 voorbeelden omwisselen. Bij de cost function voor logistic regression zou nog een kort voorbeeldje bij moeten. De uitleg over kernels bij Support Vector Machines moet duidelijker verwoord worden. Het gedeelte over distances kan verplaatst worden naar een ander hoofdstuk. Bij Clustering zou pseudocode toegevoegd mogen worden. Decision tree algorithms en Bayesian algorithms moet beter uitgelegd worden. Verder zijn er nog enkele kleinere opmerkingen. Er wordt geprobeerd om al deze feedback te verwerken voor de deadline van de eerste draft (16 mei 2016)

De actiepunten die gedaan zijn:

- Schrijven hoofdstuk IP Flows
- Afmaken Attack Classification hoofdstuk
- Herwerken van hoofdstuk "Machine learning for an IDS"
- Verder verwerken dataset Cegeka

Volgende actiepunten zijn besproken:

- Feedback verwerken van machine learning hoofdstukken
- Afmaken "implementatie" hoofdstuk
- Afmaken "evaluatie" hoofdstuk
- Afmaken "Intrusion detection systems" hoofdstuk
- Schrijven inleiding hoofdstuk
- Schrijven conclusie hoofdstuk
- Schrijven Future work hoofdstuk
- Schrijven van Nederlandstalige samenvatting
- Verder verwerken dataset Cegeka